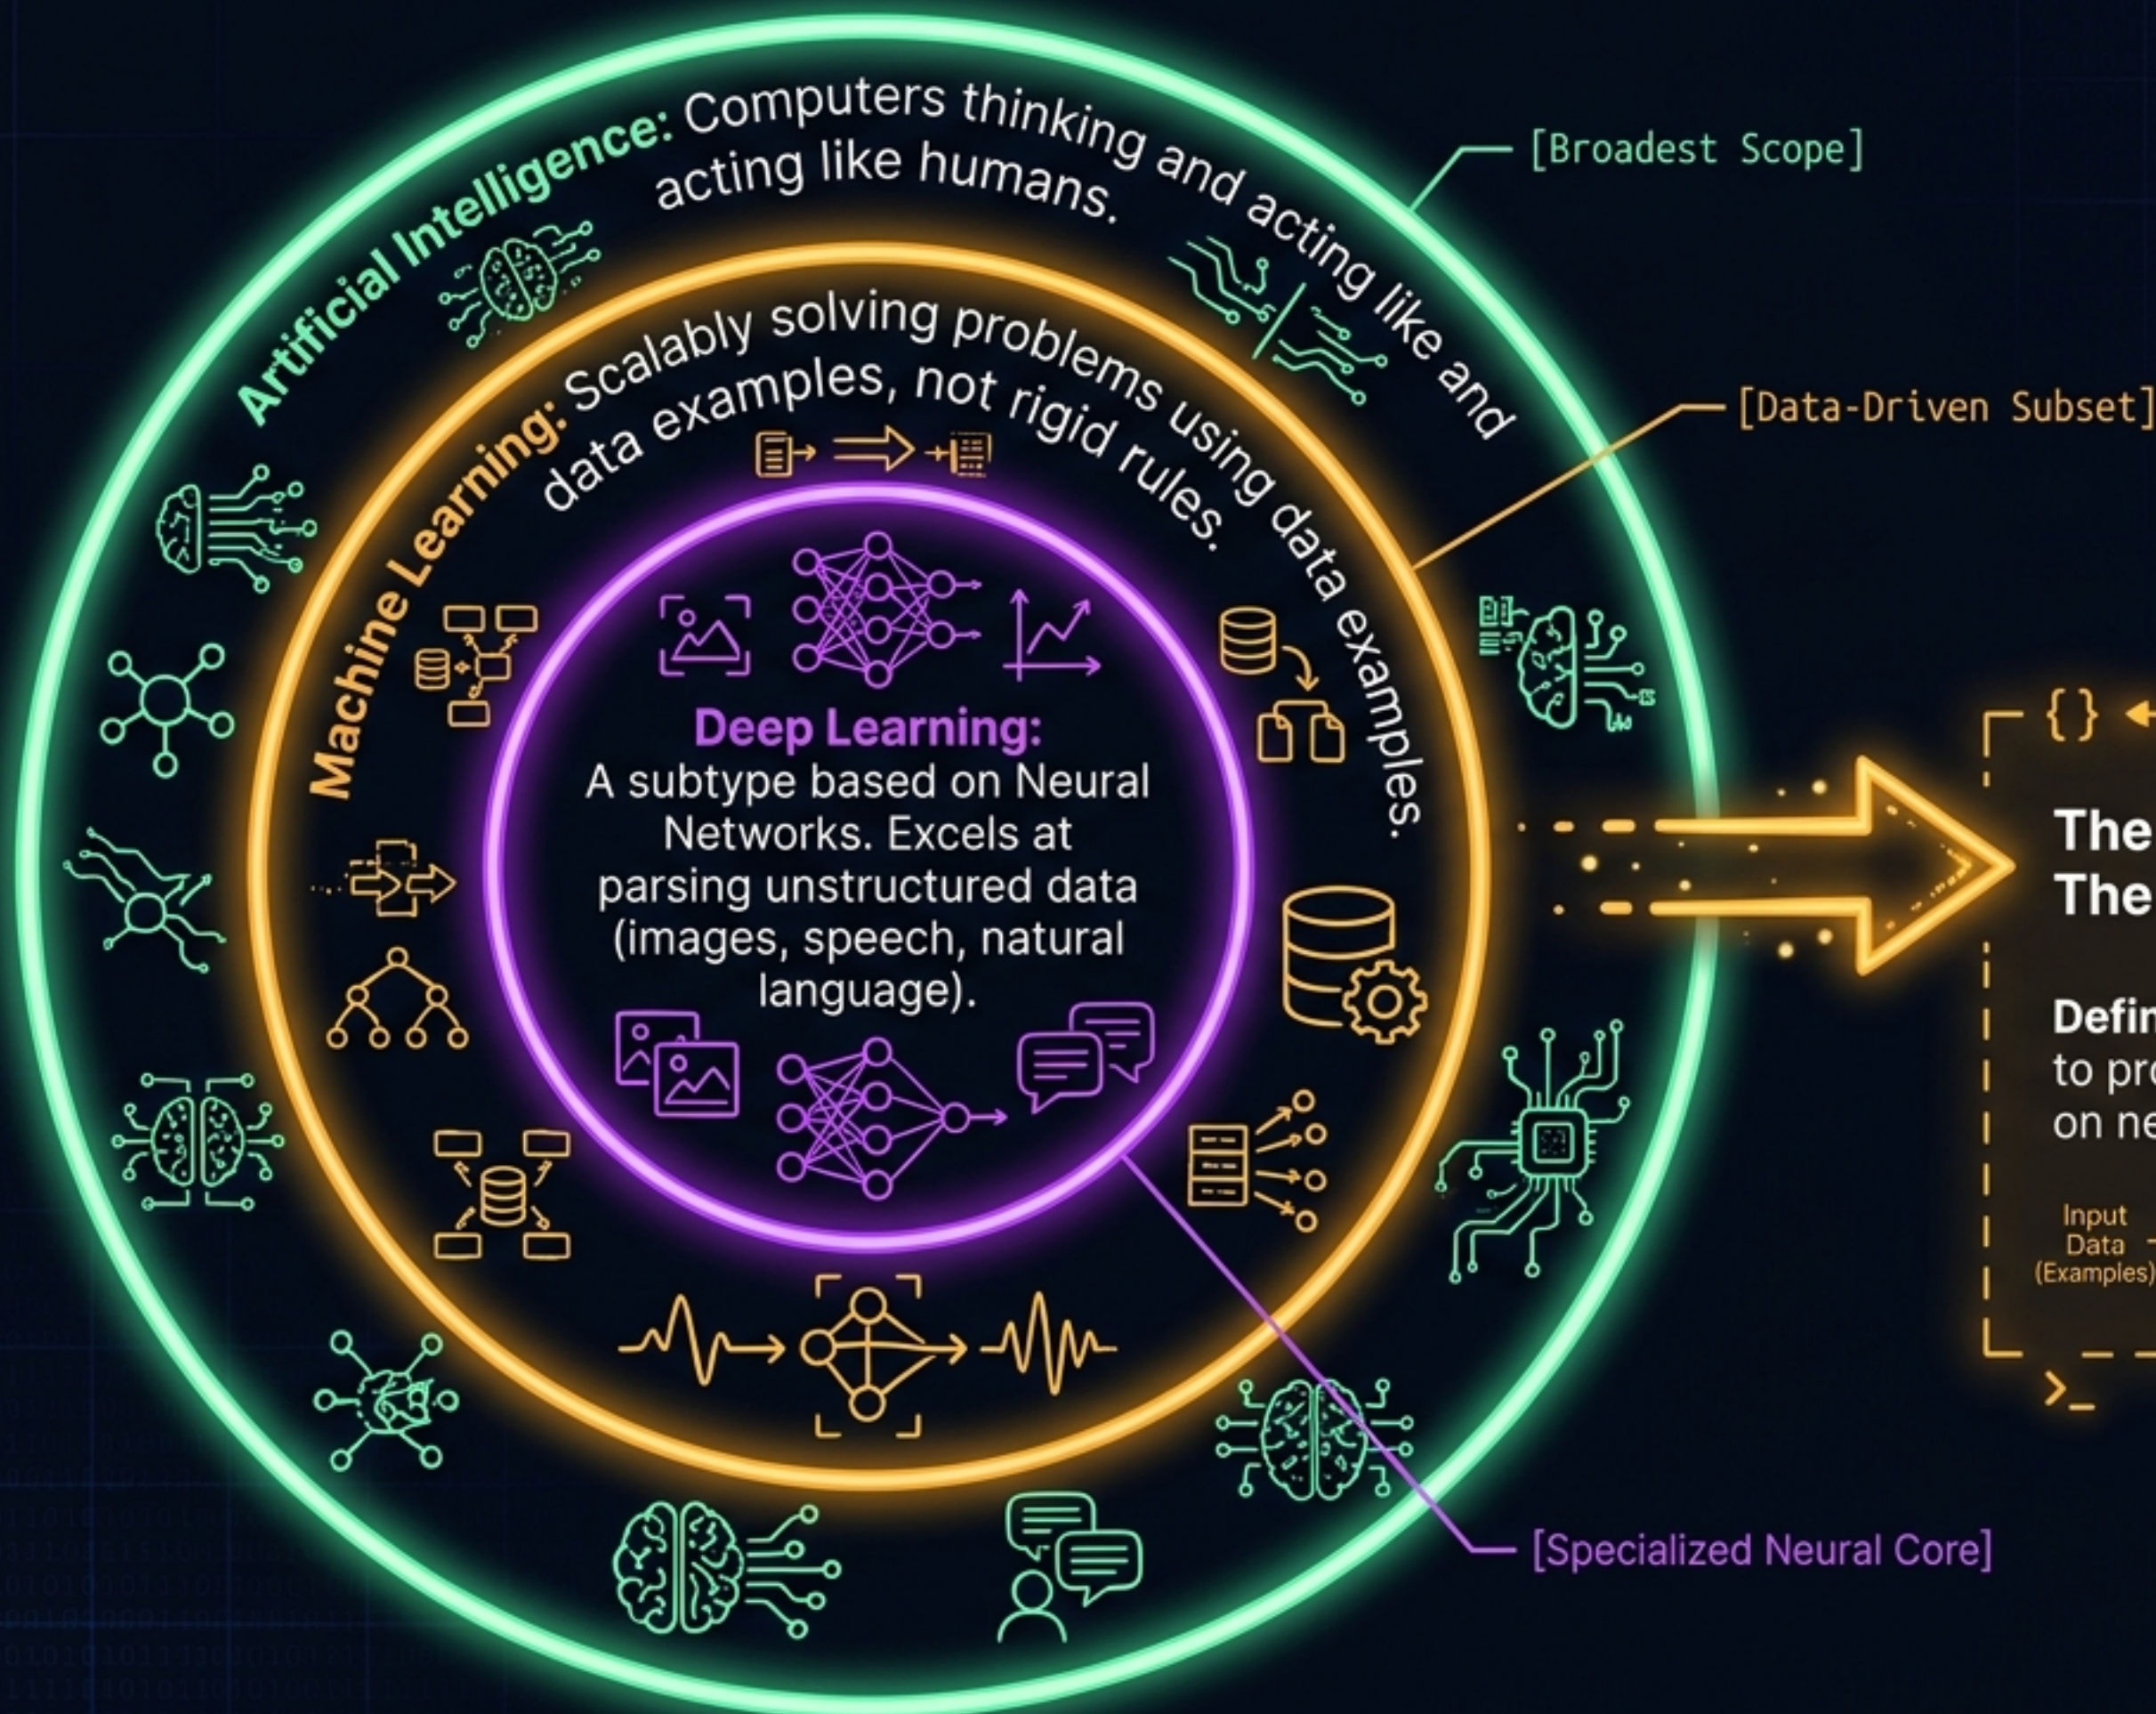


Machine Learning 101

By Michaël BETTAN



**The Core Difference:
The machine learns.**

Definition: Combining inputs to produce useful predictions on never-before-seen data.



Learning Paradigms

```
...
//kon:toDrad:
rneters.swih(-veih examplet, / aiAnt)

// salarancs:
dota.esDrae(stuases)
...
```

```
class aozu8albIn(tr => {
  if (80000 == oneuserop-trackfen) {
    restain.presaxaracter(80senou8)
    ...
  } else {
    inut prwvineer {
      R88.pemaboeCa < => beonlypsijbFlane!!
    }
    ...
  }
}
```



Supervised Learning

Learn from examples, apply labels to data.

- **Binary Classification:** Yes/No (e.g., Credit fraud)
- **Multi-class Classification:** Discrete segments (e.g., Personas)
- **Regression:** Continuous values (e.g., Forecasting spend)
- **Matrix Factorization:** Recommenders (e.g., Netflix)

Unsupervised Learning

Detect known patterns without examples.



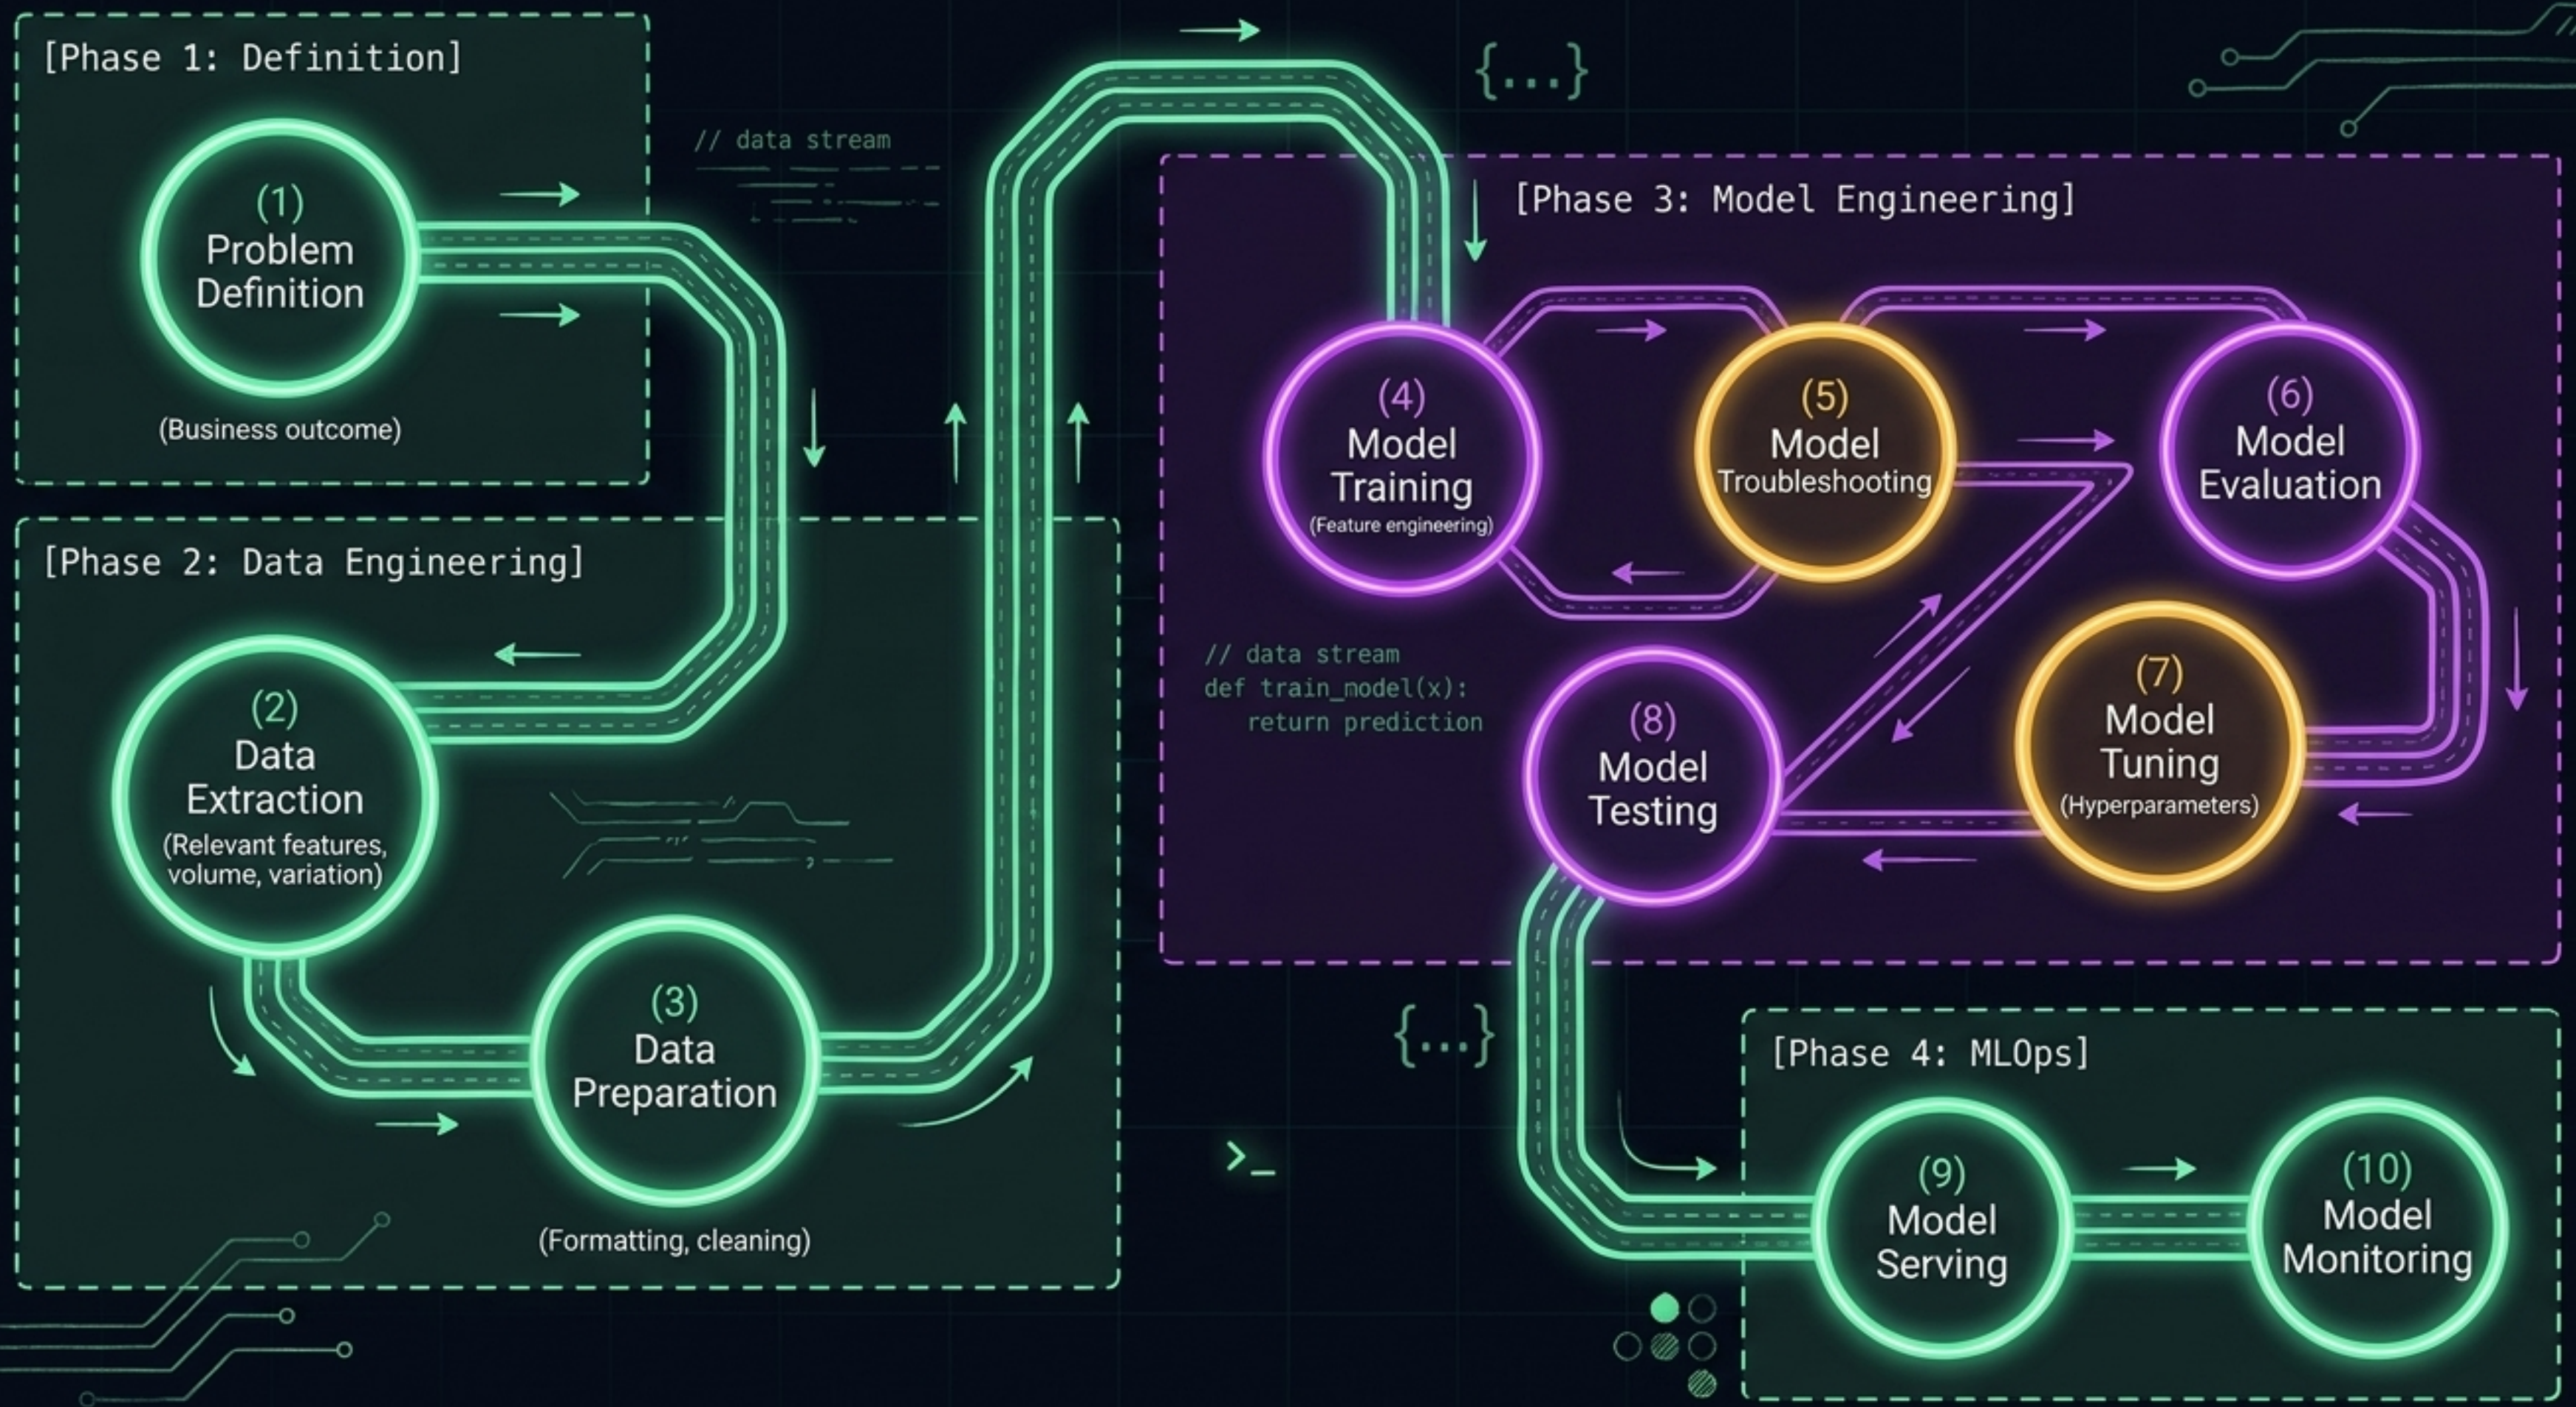
- **Clustering:** Identifying similarities in groups
- **Anomaly Detection:** Identifying abnormalities in data



Reinforcement Learning

Learn from environment via exploration & exploitation.

- **Positive/negative reinforcement** to complete a task
- **Examples:** Chess, maze navigation



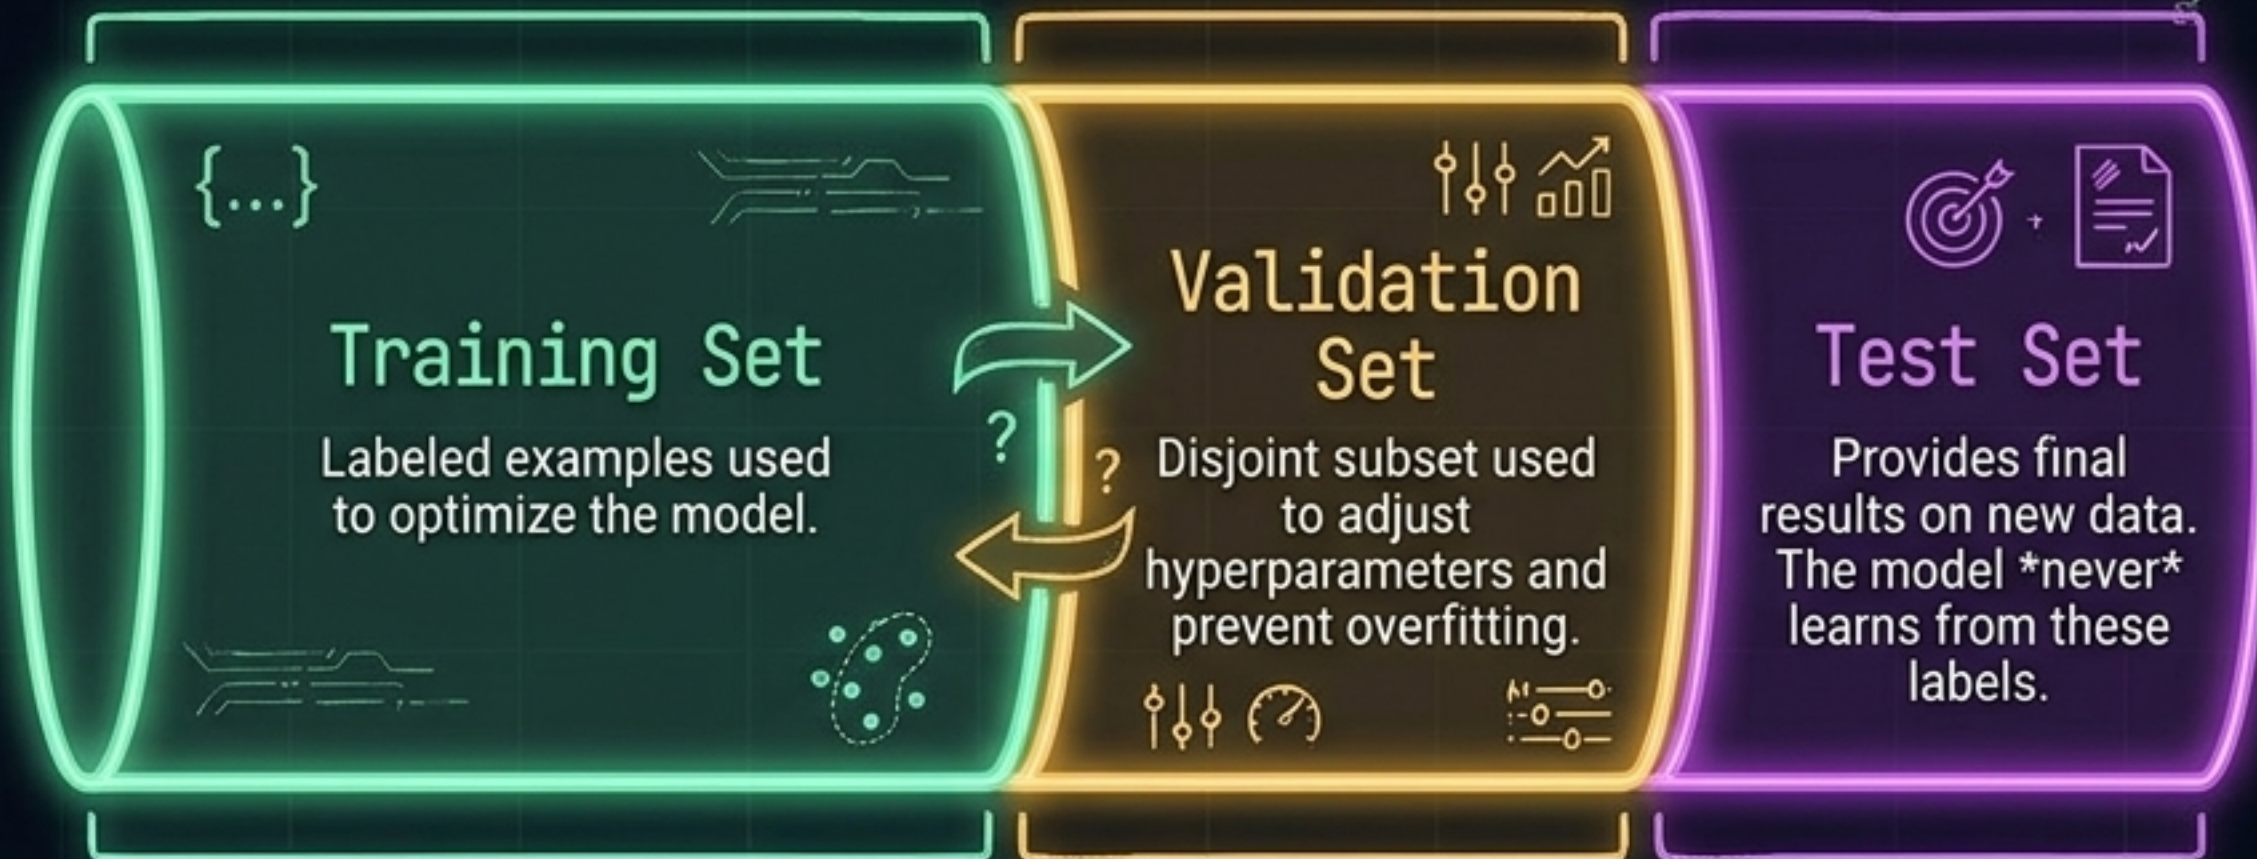
```
ftans:
// lion-todrad:
  rnoters.swin(-veih examplet, / stAnt)

// selerance:
  data.splitep(shusses)
```

// dataset split



[Raw Data Dump]



/* critical warnings */

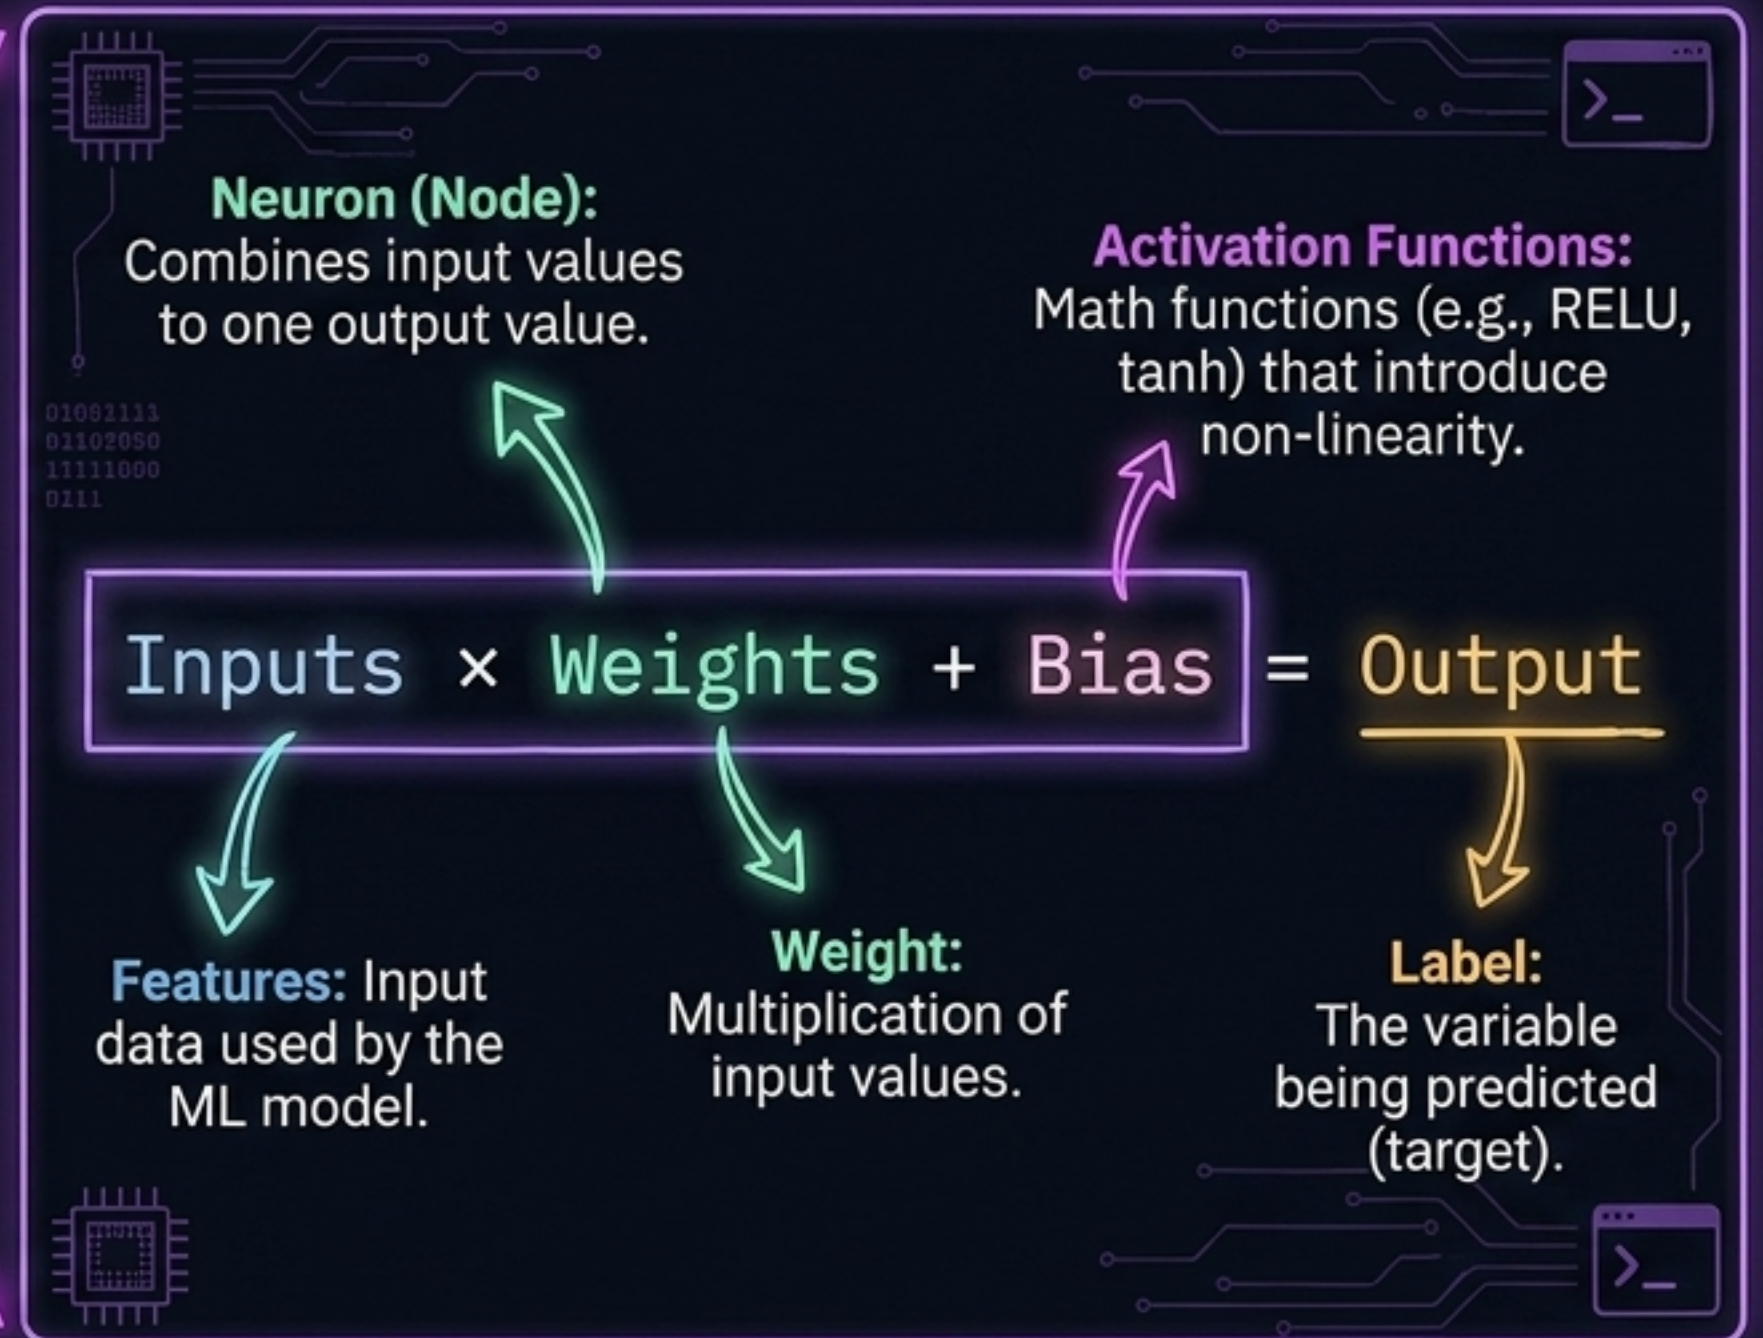
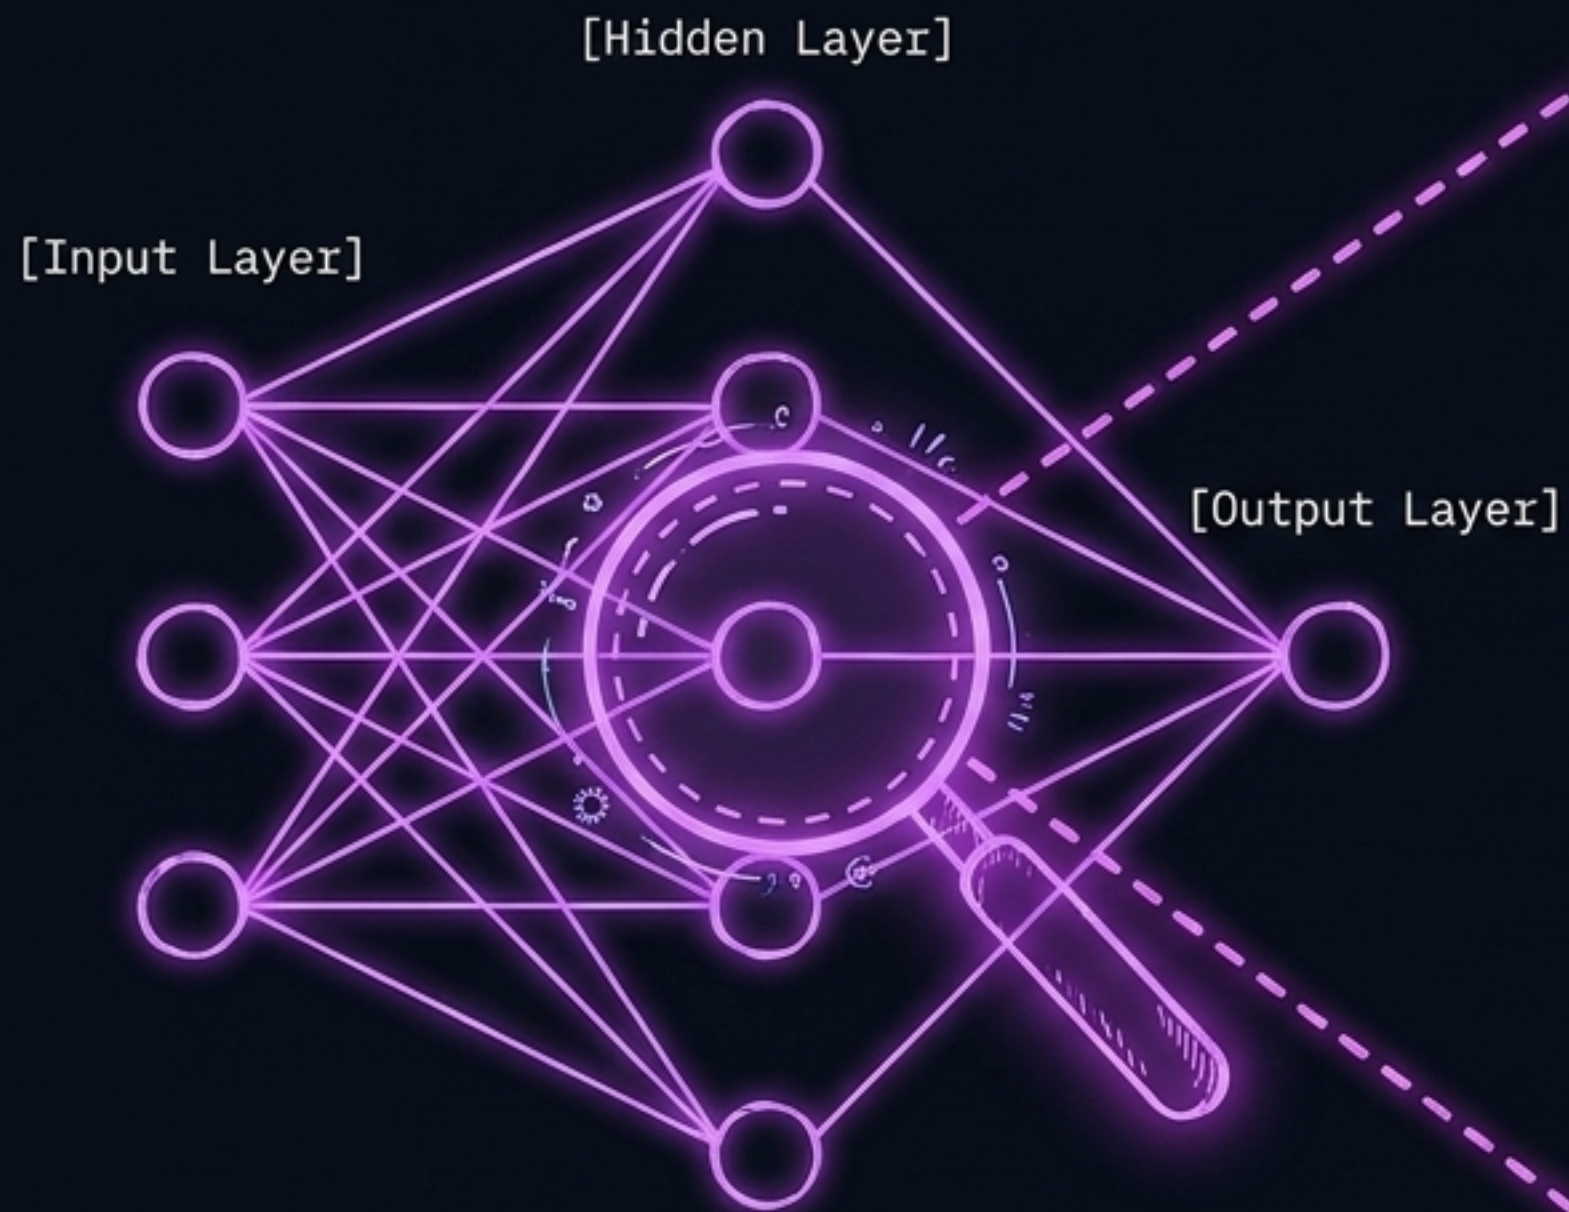
Data Leakage

Input features contain "leaked" info about the target that won't be available in production.
Causes strong testing but weak deployment.

Training-Serving Skew


Mismatch between input features available during training versus what is actually available at inference time.

Anatomy of a Neural Network



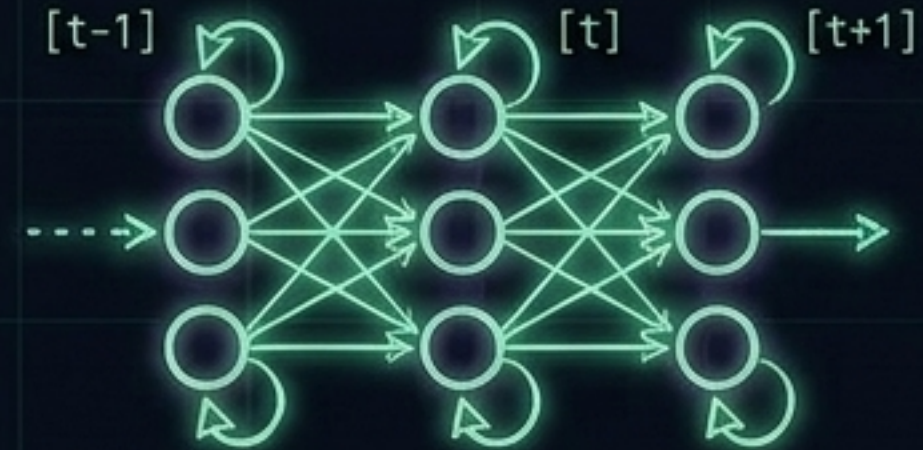
Flow Typologies

Feedforward (FFNN)



Information moves strictly forward.
No recursive connections.

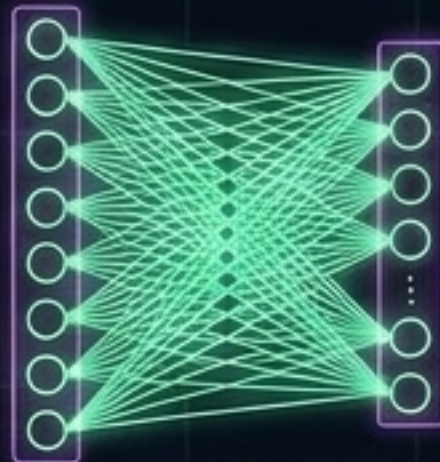
Recurrent (RNN)



Optimized for sequential data;
previous runs feed into the next.


Scale Typologies

Wide Networks



Focus on Memorization (many features).

Deep Networks



Focus on Generalization (many hidden layers).

Deep & Wide Networks: Combining both creates a perfect fit for recommendation engines.

The Bias-Variance Tradeoff

Overfitting

Good training, poor test.
High Variance.
Sensitive to small training
fluctuations.

Cause: Memorized the data,
redundant features, or noise.

High Variance

High Error Scenario



Ideal State

Low Bias, Low Variance.
Generalized perfectly.

Optimal
Generalization

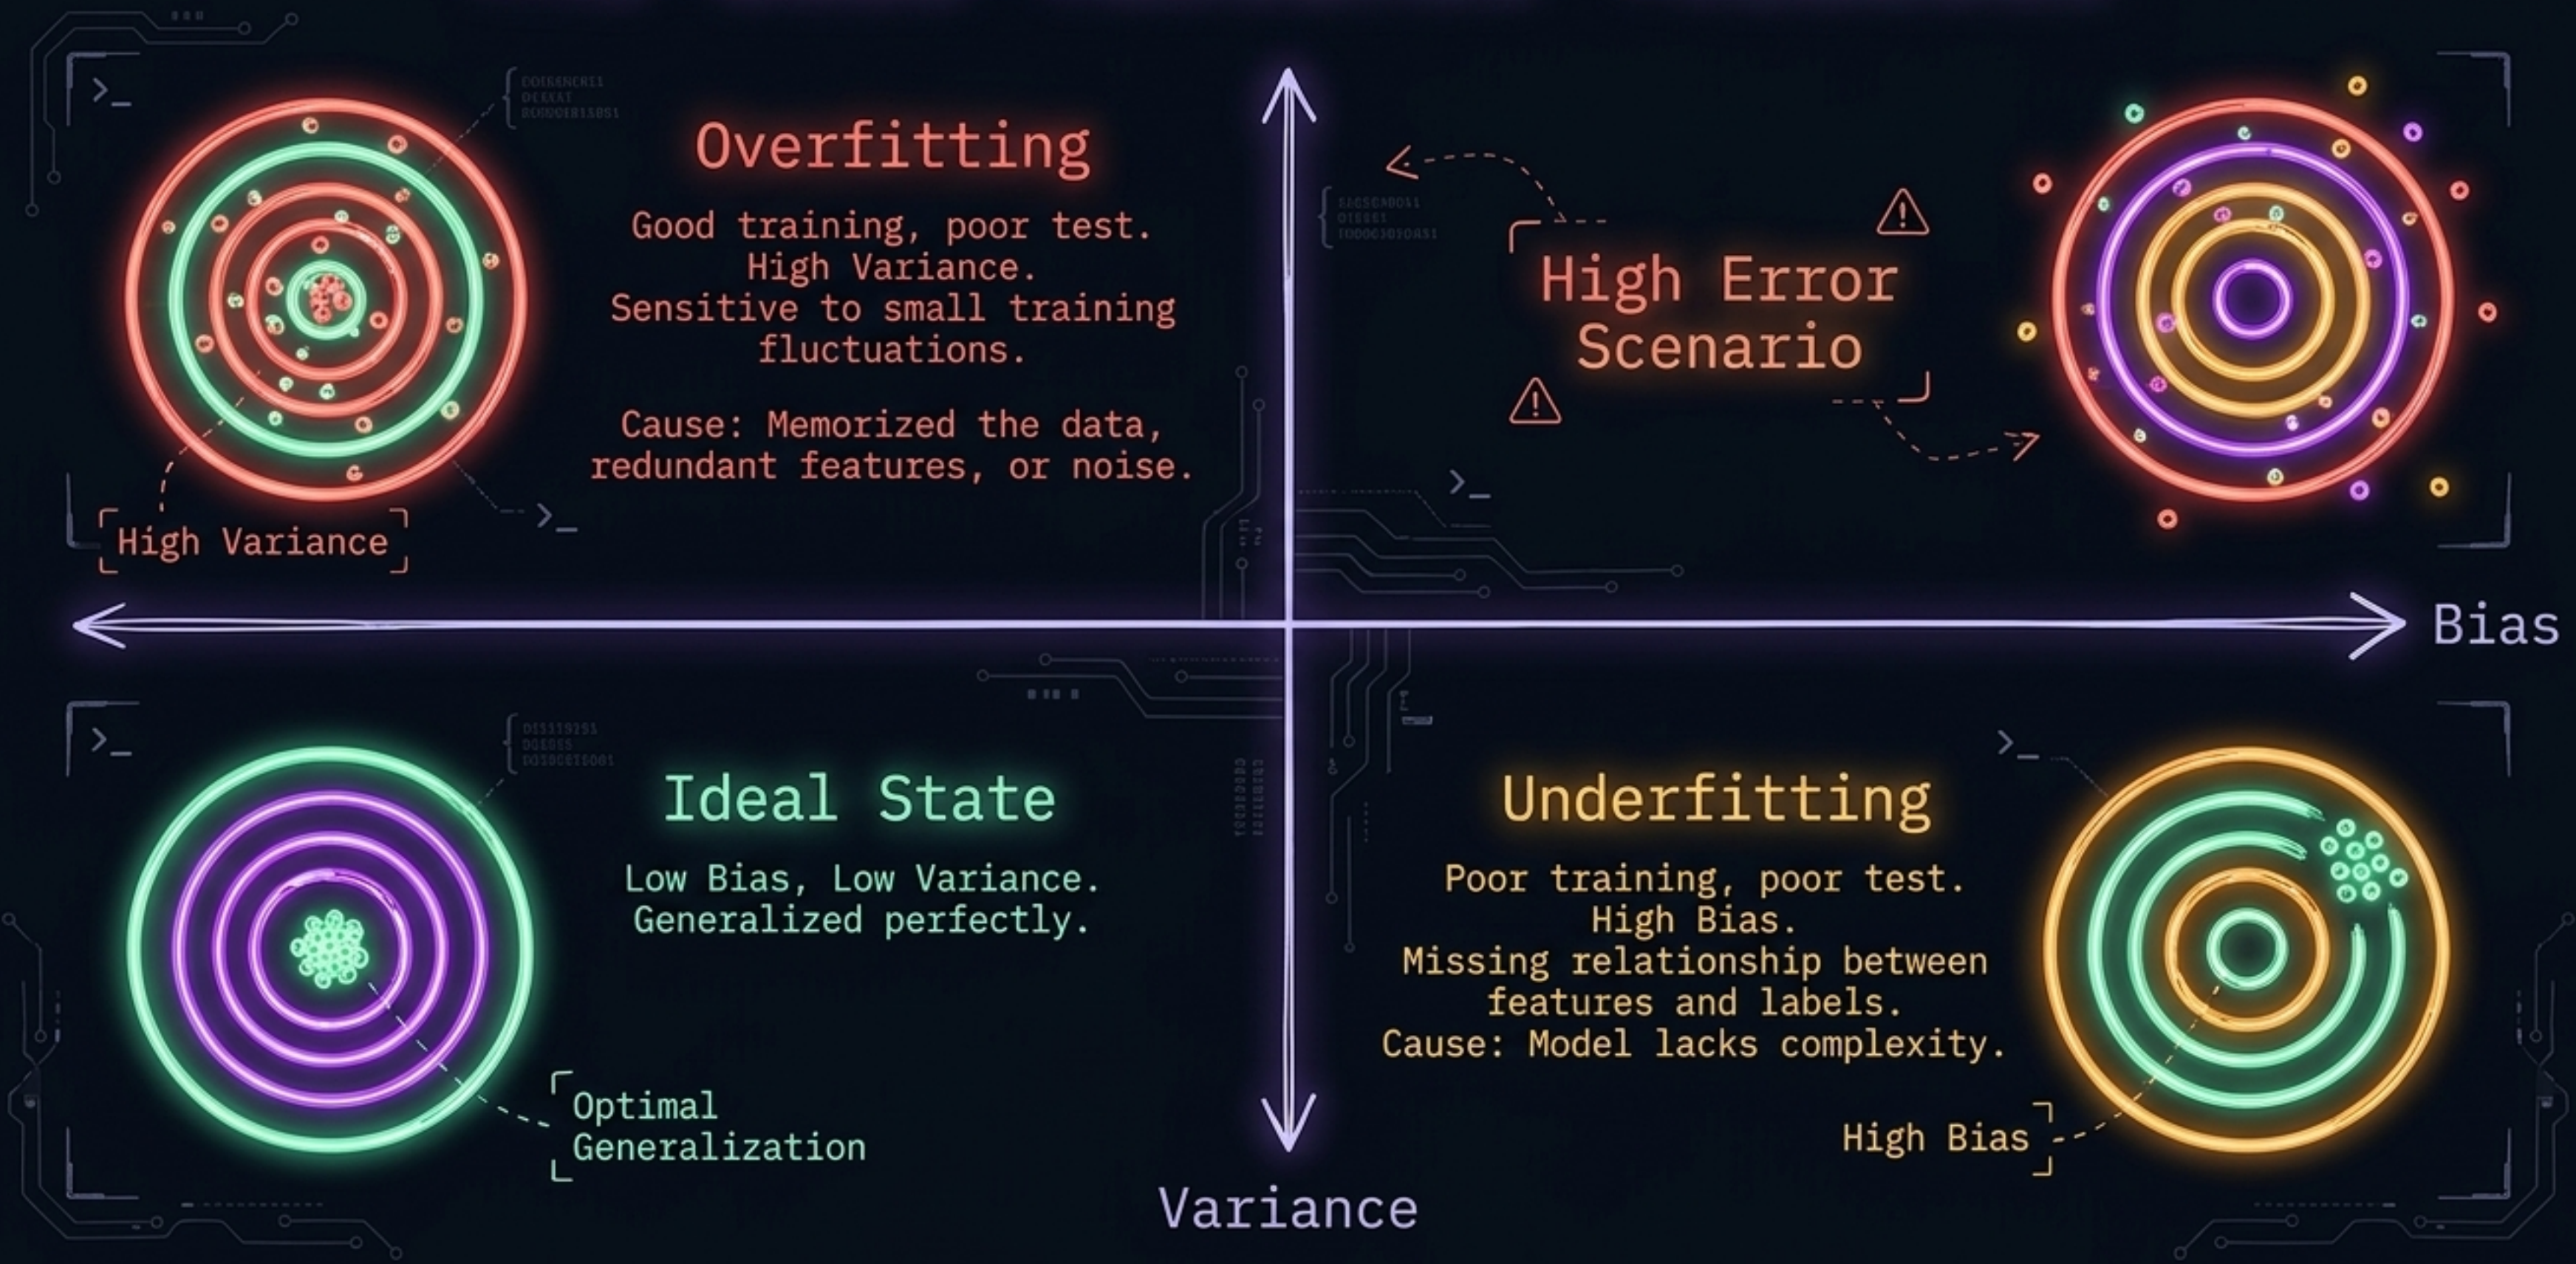
Underfitting

Poor training, poor test.
High Bias.
Missing relationship between
features and labels.
Cause: Model lacks complexity.

High Bias

Bias

Variance



Early Stopping

Halting training when validation loss starts to increase.



Dropout

Randomly dropping nodes during training to improve generalization.




Max-Norm

Limits absolute magnitude of network weights.



$$\|w\|_2^2 = \sum_i w_i^2$$




L2 Ridge

Penalizes squared weights; keeps weights approximately equal in size.

L1 Lasso

Penalizes absolute weight values; drives least useful features to 0.



Combat Overfitting (Regularization)

[CRITICAL: GENERALIZATION FOCUS]

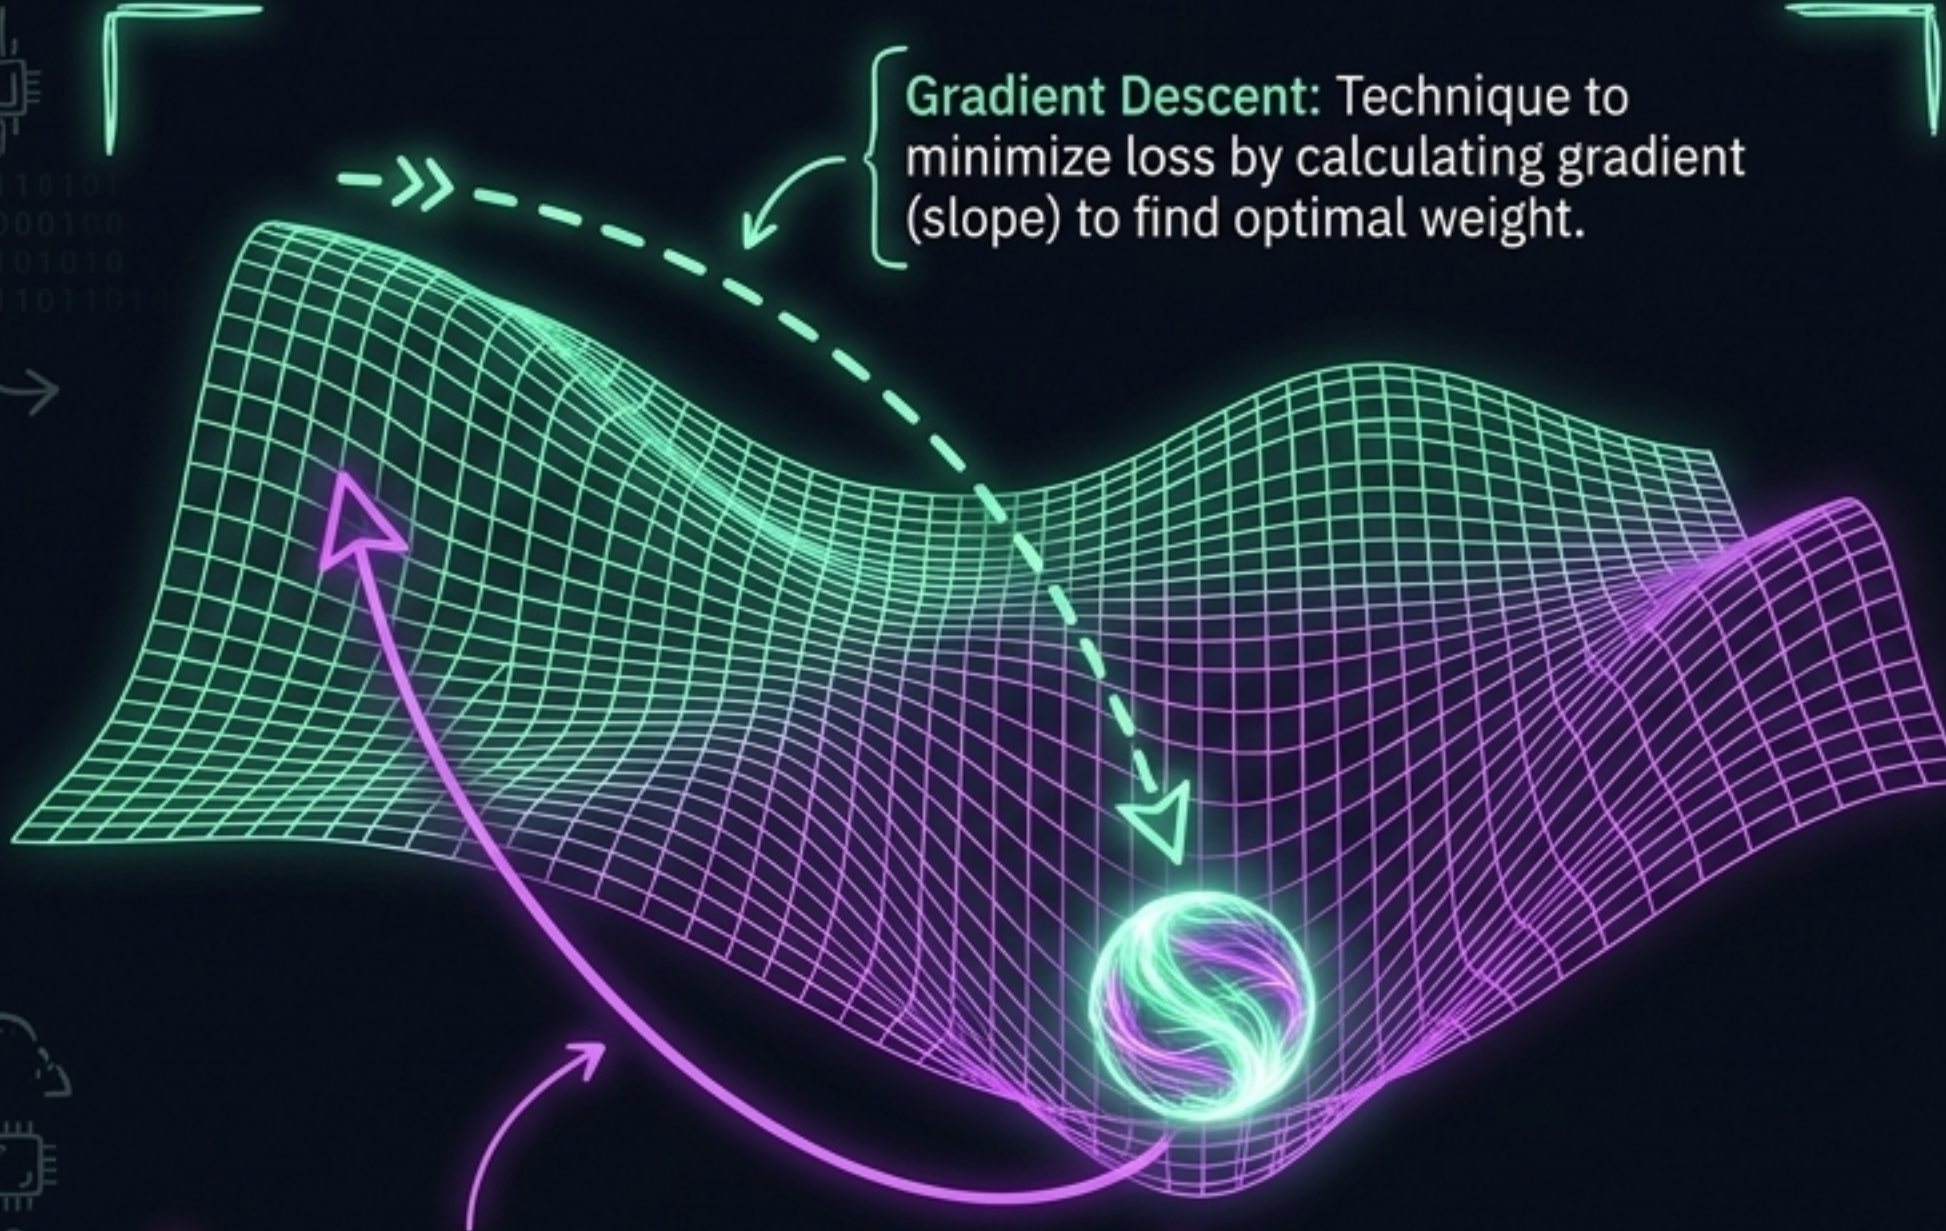


Data Augmentation

Transforming existing examples to artificially boost dataset diversity.




The Optimization Engine: Finding the Answer



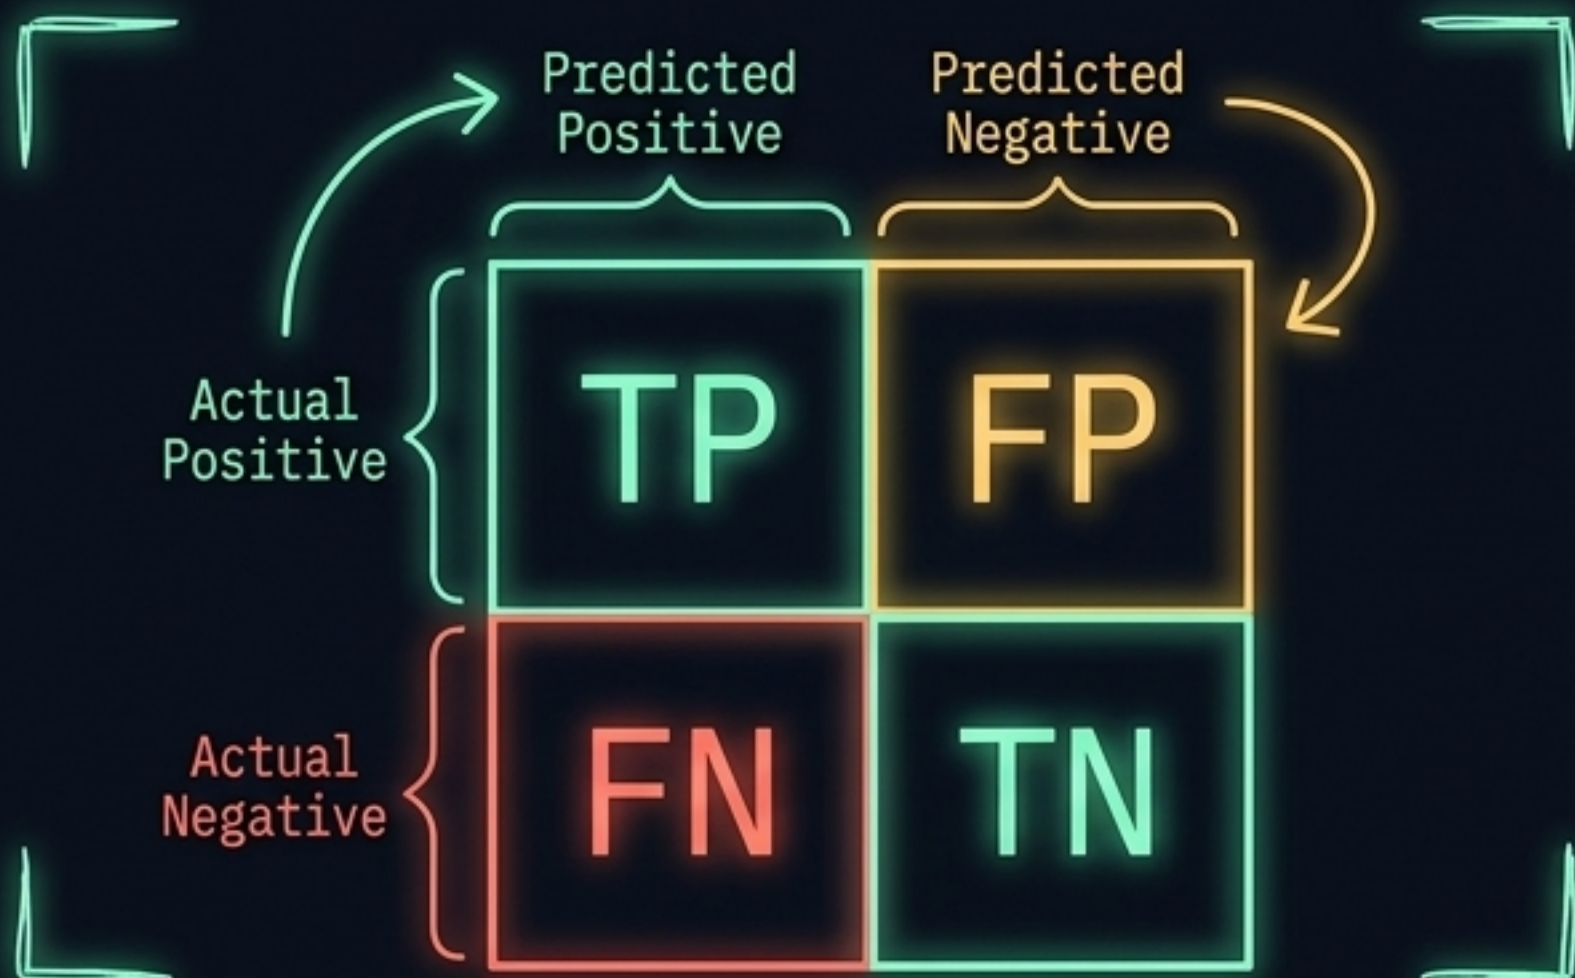
Gradient Descent: Technique to minimize loss by calculating gradient (slope) to find optimal weight.

Backpropagation: Efficient algorithm mapping input-output pairs back through the network to calculate gradient descent.

Key Engine Variables

- **"Epoch"**: A single pass through the entire training dataset.
- **"Optimizer"**: Operation (e.g., Adam, Adagrad) that changes weights/biases to reduce loss.
- **"Learning Rate"**: Speed at which optimizers adjust weights.  Risks non-convergence if too high.
- **"Converge"**: State where loss stabilizes and optimal answer is reached.
- **"Batch Gradient Descent"**: Calculates across whole dataset (slow for large data).

Classification Metrics



$$\text{Accuracy} = (\text{TP} + \text{TN}) / \text{Total}$$

$$\text{Precision} = \text{TP} / (\text{TP} + \text{FP}) \text{ [% of positive class correctly classified]}$$

$$\text{Recall} = \text{TP} / (\text{TP} + \text{FN}) \text{ [% of actual positive labels identified]}$$

$$\text{F1 Score} = 2 * (\text{Precision} * \text{Recall}) / (\text{Precision} + \text{Recall})$$

Other: AUC PR, AUC ROC, Log loss

Regression Metrics

Measuring mean squared errors for continuous values.

MAE : Mean Absolute Error

RMSE : Root Mean Squared Error


RMSLE : Root Mean Squared Logarithmic Error

Model Serving architectures





Online Serving

Synchronous & Real-Time



Quickly returns a prediction but only accepts ONE prediction request per API call.

Use Case: Models embedded in live applications dependent on quick, low-latency turnarounds.



Batch Serving

Asynchronous & Bulk



Model waits to process ALL prediction requests before returning values.

Use Case: Making many prediction requests at once where system latency is not the primary constraint.




Inference: Applying a trained model to unlabelled examples.

Disambiguation: The Three Faces of Bias

Math Bias


The intercept or offset from an origin. The 'b' in the equation $y = wx + b$. Represents the value of output given a weight of 0.



BIAS

Prediction Bias

The mathematical difference between the average of predictions and the average of actual labels in the dataset.



Ethics/Fairness Bias

Unintentional unfairness or stereotyping baked into algorithms or underlying training data. Must be explicitly tested for during monitoring.



The Human Ecosystem: MLOps Personas



Product Manager --- Defines insights, objectives, and problem definition.



Data Analyst --- Queries and analyzes raw data.



Data Engineer --- Ensures clean, useful data extraction and preparation.



Data Scientist --- Builds and trains models that work.



ML Developer --- Integrates models into intelligent applications.



ML Engineer --- Operates models in production (DevOps for ML).

1. Problem Def.

2. Data Extraction

3. Data Prep

4. Model Training

5. Model Eval

6. Model Packaging

7. Integration

8. Deployment

9. Monitoring

10. Retraining

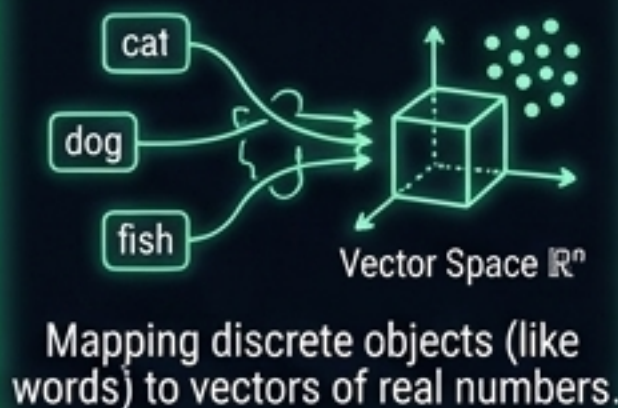
Ops Focus (Cloud Monitoring):

Alerting for traffic patterns, error rates, latency, and Data Skew (drift in data over time requiring model refresh).

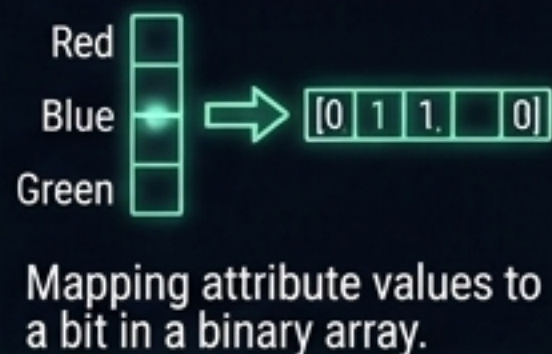
ML Developer Cheat Sheet: Key Concepts Glossary

Data Transformation


Embeddings



One-Hot Encoding



Normalization




Architecture & Scale

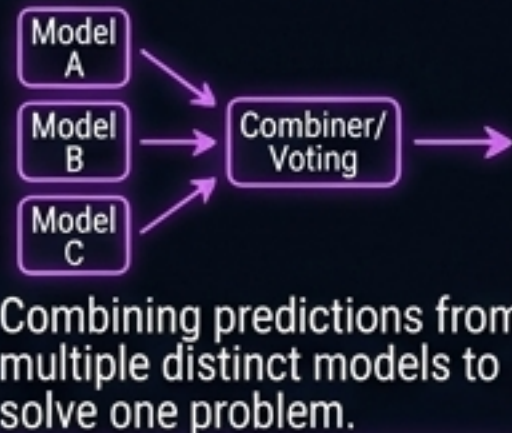
Tensor



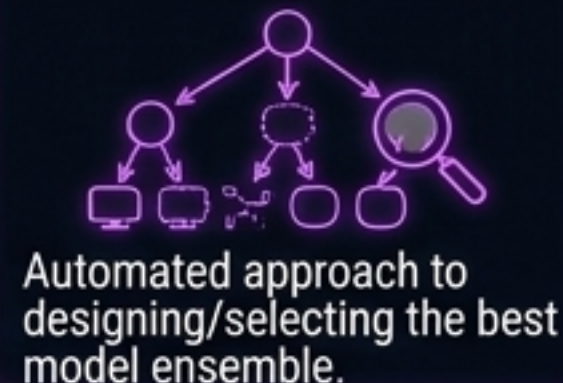
Feature Store



Ensemble Learning



Neural Architecture Search (NAS)



// Core Equation
Inference = Scoring = Predictions
(Applying a trained model to unlabelled examples)