# Prompt Engineering by Michaël Bettan

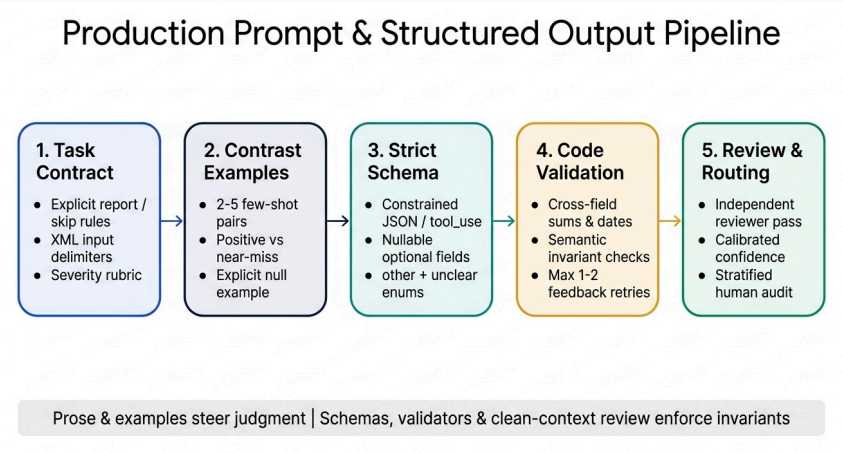
### Definition

Prompt engineering is the discipline of specifying task contracts, evaluation rubrics, input boundaries, demonstration examples, and output schemas so a language model produces reliable, verifiable results across varied production inputs—paired with programmatic validation, explicit reasoning boundaries, and staged prompt chains where prose instructions alone are insufficient.

### Why production prompting looks different from chat prompting

In an interactive chat, a human spots a vague answer and clarifies on the next turn. In an automated code reviewer, document extractor, or customer-operations pipeline, nobody sits in the loop unless your system routes the item there. Production prompt engineering therefore combines **authoring levers** (explicit categorical criteria, contrast examples, input delimiters, normalization rules) with **API enforcement levers** (constrained JSON schemas, strict tool definitions, `tool_choice`, adaptive reasoning effort, bounded validation retries, and multi-instance review).

> MENTAL MODEL
> Prose steers *judgment*; schemas and code enforce *invariants*. Use prompt instructions and few-shot examples to define ambiguous boundaries and domain semantics, and use API schemas, validators, and fresh reviewer contexts to guarantee syntax, catch cross-field math errors, and eliminate self-review blind spots.




- **Scope of this guide:** prompt anatomy and attention placement, explicit criteria and false-positive control, multishot design, reasoning effort vs. sampling parameters, task decomposition and iterative refinement, schema-enforced structured outputs, validation-retry and reviewer-feedback loops, batch transport selection and chained batch stages, chain-of-thought isolation and independent critique passes, caching, injection defense, and LLMOps.
- **Vendors covered:** Anthropic Claude (Messages API, `tool_use`, `output_config`, Adaptive Thinking, Message Batches) and Google (Gemini API, Vertex AI `response_schema`, `FunctionCallingConfig`, Thinking levels, Batch Prediction, Vertex AI Prompt Optimizer).

## Prompt Anatomy, Roles & Attention Placement

### Message roles and trust hierarchy

| Channel | Claude Messages API | Gemini / Vertex AI | What belongs here |
|---|---|---|---|
| **System instruction** | Top-level `system` parameter | `system_instruction` in config | Role identity, non-negotiable behavioral rules, categorical criteria, global output conventions |
| **Reference & examples** | Early `user` content blocks or `system` blocks (with `cache_control`) | Cached prefix contents or early `user` parts | Domain glossary, schema normalization rules, `<examples>`, large source documents |
| **Untrusted payload** | Delimited block inside `user` (`<document>…</document>`) | Delimited part inside `user` | Raw customer messages, pull-request diffs, OCR text, web pages |
| **Task trigger** | Final lines of the `user` turn | Final part of the `user` turn | The specific question, target action, and reminder of edge-case rules |

### Structural delimiters and edge placement

- **Separate instructions from data:** wrap variable inputs in unambiguous tags such as `<document>`, `<diff>`, or `<customer_ticket>` rather than loose quotation marks. Clear boundaries stop multi-paragraph inputs from bleeding into your rules and make indirect prompt-injection instructions easier for the model to recognize as data.
- **Respect the U-shaped attention curve ("Lost in the Middle"):** across long prompts, recall is highest at the opening and closing edges of the window and weakest in the middle interior. Place large reference documents near the top (where they also form a stable cacheable prefix), put key findings summaries ahead of raw appendices, label sections with explicit headers, and state the exact task instructions and critical constraints at the very end.
- **Positive framing vs. hard negative boundaries:** frame style and workflow instructions positively (*"Write a 3-bullet summary citing section numbers"* outperforms *"Don't write long paragraphs or omit citations"*). Reserve strict negative directives (`NEVER`, `DO NOT`, `ONLY`) for safety gates, scope exclusions, and anti-fabrication rules (*"Return `null` when a date is absent; never estimate or infer missing dates"*).
- **Authoring checklist (Task · Context · Examples · Evaluate · Iterate):** define a concrete verb, role, and target audience (**Task**); supply the minimum situational facts and source boundaries needed (**Context**); attach 2–5 representative demonstrations (**Examples**); grade outputs on a labeled test set (**Evaluate**); and refine one variable per run (**Iterate**).

### XML tag disambiguation rules

Claude has no reserved tag vocabulary: any descriptive tag name works, provided it is used consistently. Gemini prompts benefit from the same discipline (clear prefixes or tags that separate instructions, context, and input).

| Rule | What goes wrong without it | Practice |
|---|---|---|
| **One tag per content type** | Instructions, reference text, and live input blur together | Wrap each kind separately: `<instructions>`, `<context>`, `<examples>`, `<input>` |
| **Nest to express hierarchy** | The model cannot tell which passage a rule refers to | `<documents>` containing `<document index="2">` with inner `<source>` and `<content>` tags |
| **Refer to tags by name** | *"Use the document above"* is ambiguous once several blocks exist | *"Answer only from `<contract>`; apply tone rules from `<style_guide>`"* |
| **Keep names stable** | Renaming tags between prompt versions or examples confuses the model and breaks cached prefixes | Freeze tag names in the versioned template; change content, not structure |
| **Neutralize tag look-alikes in payloads** | Untrusted text containing `</input>` can fake the end of the data block | Escape or strip tag-like strings from the payload before insertion |
| **Name the output container** | Free-form replies mix commentary with the result | Ask for the result inside `<answer>` for human-facing text; use a JSON schema for machine consumers |

**RETIRED PATTERN — ASSISTANT PREFILLING**

Older guides recommended ending a request with a partial `assistant` turn such as `{"role": "assistant", "content": "{"}` to force JSON output. Modern frontier models (Claude 4.6 and later) **reject assistant prefilling with HTTP 400**. Enforce structure with native JSON Schema decoding (`output_config.format` / `response_schema`) or strict tool calls instead.

## Explicit Criteria, Rubrics & False-Positive Control

### Why vague instructions and confidence prompts fail

In automated code review, compliance triage, and defect detection, **precision** (the share of flagged items that are genuinely actionable) governs whether engineers read or mute your system. Telling a model to *"be conservative"*, *"only flag high-confidence issues"*, or *"check that comments are accurate"* fails because it never defines the decision boundary—the model's internal threshold stays uncalibrated while noisy edge cases still slip through.

| Weak / vague phrasing | Failure mode in production | Explicit categorical specification |
|---|---|---|
| *"Check whether code comments are accurate"* | Flags minor wording preferences, TODOs, and abbreviated docstrings | *"Report a comment defect only when a comment's stated behavior directly contradicts the executable code path"* |
| *"Only report high-confidence bugs; be conservative"* | Model still emits stylistic nitpicks labeled `"high confidence"` | **Report:** null dereferences, auth bypasses, resource leaks, SQL/command injection. **Skip:** naming conventions, formatting, and idioms already established in the repository |
| *"Assign a severity of low, medium, or high"* | Severity labels drift across runs and files | Define each tier with explicit impact rules **plus one concrete code example per severity level** |

### Protecting developer trust in automated pipelines

- **One noisy category poisons every category:** if your PR reviewer is 92% accurate on security bugs but 35% accurate on "maintainability" warnings, developers stop reading *all* comments from the bot within days.

- **Measure precision per category, not only in aggregate:** track acceptance and dismissal rates separately for each rule family (`security`, `concurrency`, `null-safety`, `comment-drift`).
- **Disable noisy categories immediately while fixing them:** when a category produces excessive false positives, **turn that category off** in production right away to restore trust in the high-precision categories, then rewrite its criteria and few-shot contrast examples offline against an evaluation set before re-enabling it.

| Policy component | Role in the prompt contract | Concrete specification |
|---|---|---|
| **Report categories** | Enumerates the only defect classes eligible for comments | `defect` (runtime exceptions, off-by-one bounds, lost updates, unhandled promises); `security` (broken access checks, unsanitized shell/SQL/HTML sinks, secret exposure); `comment_contradiction` (a guarantee such as *"thread-safe"* or *"returns UTC"* violated by the code) |
| **Skip categories** | Explicitly suppresses high-noise style feedback | Formatting, variable naming, import ordering, missing docstrings, and established local repository idioms |
| **Severity rubric** | Anchors each severity tier to concrete runtime impact | `critical`: exploitable security flaw or guaranteed crash/data loss on valid input (e.g., `if is_admin or True:`); `major`: wrong business result or unhandled edge case on reachable inputs (e.g., dividing by `len(items)` before checking emptiness) |

EXAM TRAP

When a scenario shows developers ignoring an automated PR reviewer due to false positives, avoid answers that *"lower temperature to 0"*, *"tell the model to only report findings above 90% confidence"*, or *"keep all categories live while adding a disclaimer"*. Temporarily disable the noisy category and replace vague prompts with explicit report/skip criteria and severity examples.

## Few-Shot & Multishot Design

### When few-shot examples beat longer instructions

When prose rules are already clear yet outputs still vary in structure, tone, or boundary judgment, **2–5 targeted few-shot examples** (wrapped in `<examples><example>…</example></examples>`) are the highest-leverage intervention. Demonstrations teach **generalizable judgment** on novel patterns rather than brittle keyword matching on enumerated cases.

| Engineering goal | How to construct the `<example>` set |
|---|---|
| **Consistent output format** | Every example shows the complete target shape (for a ticket classifier: `intent`, `urgency`, `evidence_span`, `reply_draft`) with identical key order and casing |
| **False-positive suppression** | Use **contrast pairs**: one example showing a real defect that *is* reported alongside a plausible look-alike or valid repo idiom that produces `[]` (no finding) |
| **Ambiguous routing / tool choice** | Include a brief `<reasoning>` trace inside the example explaining *why* action A beats plausible alternative B |
| **Varied document extraction** | Demonstrate extraction across structural layouts (a prose paragraph with inline citations, a formal bibliography table, informal units like `"half a dozen"`, and a document missing an optional field returning `null`) |

### Structuring the example block

Keep demonstrations physically and semantically separate from the live request so the model never mistakes a sample input for the one it must process:

```
<examples>
  <example>
    <input>Card charged twice for order 5521</input>
    <ideal_output>{"intent": "billing", "urgency": "high"}
</ideal_output>
  </example>
</examples>
<input>{{TICKET_TEXT}}</input>
```

- **Label the role of examples:** state that examples illustrate format and judgment only, and that names, amounts, and wording must come from the live `<input>`. Unlabeled examples get copied verbatim.
- **Show the target pattern, not the anti-pattern:** Gemini's prompt design guidance notes that examples of what to do steer better than examples of what to avoid. Contrast pairs still follow this rule: the look-alike input is paired with its *correct* (empty or benign) output.
- **Keep formatting identical across examples:** same tags, same key order, same whitespace conventions. Inconsistent demos teach inconsistent output.
- **Place examples in the static prefix:** examples before the live input stay cacheable and leave the task trigger at the end of the prompt.

### Preventing few-shot overfitting and bias

- **Randomize class order (class mixing):** listing four `"escalate"` examples followed by one `"resolve"` example biases the model toward escalating borderline inputs. Interleave outcomes and categories across examples.
- **Vary surface structure:** if every example invoice uses `USD` and three line items, the model may copy `USD` on a `EUR` document or drop a fourth line item. Vary lengths, field presence, and phrasings while keeping the schema invariant.
- **Demonstrate `null` explicitly:** to cure hallucination on missing required-looking fields, include at least one example where the source omits a field and the target output sets `"renewal_date": null` rather than guessing.

| Sample `<input_doc>` snippet | Target structured extraction | Boundary taught by the example |
|---|---|---|
| *"Memo (14 May 2026): Pilot renewal approved for 18 months; PO number pending finance sign-off."* | `effective_date`: `"2026-05-14"`, `term_months`: `18`, `po_number`: `null`, `status`: `"approved"` | Normalizes informal prose dates to ISO-8601 while setting an explicitly pending field (`po_number`) to `null` |
| *"Verbal discussion notes — terms still under negotiation between legal teams; no dates agreed."* | `effective_date`: `null`, `term_months`: `null`, `po_number`: `null`, `status`: `"unclear"` | Demonstrates full `null` extraction plus the `"unclear"` escape hatch when no binding agreement exists |

## Reasoning, Effort & Sampling Controls

### Adaptive thinking and reasoning depth

Modern frontier models allocate variable test-time compute ("System 2" reasoning) before emitting visible output or between tool calls.

| Capability | Claude 4.6 / 4.7+ (`Messages API`) | Google Gemini (`Gemini API` / `Vertex AI`) |
|---|---|---|
| **Enable reasoning** | `thinking={"type": "adaptive"}` (interleaves automatically across tool turns) | `thinking_config={"thinking_level": "HIGH"}` (model-dependent levels: `MINIMAL`, `LOW`, `MEDIUM`, `HIGH`) |
| **Control compute budget** | `output_config={"effort": "low" \| "medium" \| "high" \| "xhigh" \| "max"}` | `thinking_level` (or `thinking_budget` on earlier Gemini 2.5 models) |
| **Legacy trap** | Manual `{"type": "enabled", "budget_tokens": N}` is deprecated on 4.6 and **returns HTTP 400 on Claude 4.7+** | Mixing legacy budget fields with level-based models fails validation |
| **Cost & latency rule** | Full internal thinking tokens are billed as output tokens even when `display: "omitted"`; use `low` effort for simple routers/classifiers | Thinking tokens count toward output billing and latency; dial down `thinking_level` on high-throughput classification |

### Sampling parameters and modern API constraints

- **Claude 4.7+ rejects non-default sampling parameters:** passing non-default `temperature`, `top_p`, or `top_k` on Claude 4.7+ returns an **HTTP 400** (and on earlier Claude 4.x models, `temperature` and `top_k` are incompatible with thinking). Frequency and presence penalties do not exist on the Claude API. Steer determinism and creativity through explicit criteria, few-shot examples, structured schemas, and `effort`.
- **Models that expose sampling (`Gemini`, open weights, older endpoints):** lower `temperature` (`0.0-0.2`) narrows token selection toward higher-probability paths for extraction and classification, while higher values (`0.7-1.0`) broaden diversity for brainstorming. Adjust `temperature` *or* `top_p`, rarely both simultaneously.
- **Low temperature is never a determinism guarantee:** floating-point batching, MoE routing, and hardware concurrency still produce run-to-run variation at `temperature=0`. Grade outputs on parsed schema properties and invariants, never exact raw string equality.
- **`max_tokens` / `max_output_tokens` is a hard cutoff, not a length instruction:** hitting the ceiling stops generation mid-token (`stop_reason == "max_tokens"` on Claude, `FinishReason.MAX_TOKENS` on Gemini), producing truncated, unparseable JSON. Ask for conciseness in the prompt/schema while sizing `max_tokens` with generous headroom.

### Reasoning prompt patterns (when explicit scaffolding helps)

- **Chain-of-Thought (CoT) in structured tasks:** when thinking mode is unavailable or disabled for latency, add a concise `"reasoning"` or `"evidence_quote"` string field *ahead* of classification fields in your schema so the model grounds its decision before committing to a label.
- **Contrastive examples:** show both a valid deduction and a common miscalculation (such as confusing fiscal-year dates with calendar-year dates) so the boundary is explicit.
- **Step-Back prompting:** ask the model to identify the governing policy principle or physical law first, then apply it to the specific case—useful when dense surface details obscure the rule.
- **Self-Consistency (majority voting):** sample multiple independent reasoning paths on hard, verifiable problems and aggregate the final structured answer—effective offline, too costly for low-latency chat.

## Decomposition & Prompt Chaining vs. Mega-Prompts

### Choosing the right decomposition strategy

Packing extraction, cross-reference verification, policy evaluation, and executive formatting into one giant "mega-prompt" dilutes attention across competing instructions and makes failures impossible to localize.

| Pattern | When to pick it | Concrete workflow example |
|---|---|---|
| **Fixed sequential pipeline (prompt chaining)** | Stages and dependencies are known upfront; predictable multi-step transformation | **Step 1:** Extract normalized claims + dates per file → **Step 2:** Cross-check claims across files → **Step 3:** Render final compliance table |
| **Dynamic / adaptive decomposition** | Open-ended investigation where step `k+1` depends on what step `k` uncovers | Legacy test-suite expansion: map module graph → rank high-risk untested modules → inspect dependencies → adapt test plan as hidden couplings emerge |
| **Parallel fan-out + synthesis** | Independent inputs that need uniform local depth without cross-contaminating context | Summarize 25 customer-interview transcripts concurrently with one shared rubric, then run 1 synthesis call over their structured outputs |

### Iterative refinement patterns for engineering workflows

1. **Concrete input/output pairs over prose rewrites:** when prose instructions for a data transformation or formatting rule are interpreted inconsistently across edge cases, stop rephrasing adjectives—provide **2–3 concrete before/after input-output pairs** covering the tricky edge cases (such as `null` timestamps or nested arrays).
2. **Test-driven iteration:** write executable unit and edge-case tests first, run the model's implementation, and feed the exact failing test assertions back so each turn closes a verified gap.
3. **The interview pattern (clarify before building):** when designing in an unfamiliar domain or underspecified problem space, instruct the model to **ask targeted architectural questions first** (covering failure modes, cache invalidation, concurrency, state boundaries) before writing code or specs.

4. **Bundled vs. sequential bug feedback:** send **interacting** defects together in a single message so the fix accounts for shared state; fix **independent** defects sequentially so instructions stay focused.

## Structured Outputs & Schema Design

### Three ways to get structured data from a model

| Mechanism | Claude implementation | Gemini / Vertex AI implementation | Syntax guaranteed? | Best fit |
|---|---|---|---|---|
| **1. Prompt-only formatting** | System/user prose or output style | System/user prose instructions | ❌ No | Streaming human-readable Markdown, YAML plans, UI prose |
| **2. Native JSON Schema output** | `output_config={"format": {"type": "json_schema", "schema": …}}` (or `claude -p --output-format json --json-schema`) | `GenerateContentConfig(response_mime_type="application/json", response_schema=…)` | ✅ Yes (on normal completion) | Single-step document extraction, classification, CI JSON gates |
| **3. Tool-use / function-call extraction** | `tools=[{…, "strict": True, "input_schema": …}]` + `tool_choice` | `tools=[FunctionDeclaration(…)]` + `FunctionCallingConfig` | ✅ Yes (with strict mode) | Multi-schema routing (`tool_choice: "any"`), agent pipelines, mandatory first-step extraction |

### Writing a schema-constrained output contract

A schema is part of the prompt: property names, enums, and `description` strings are all read by the model while constrained decoding enforces the shape.

```
{"type": "object", "additionalProperties": false,
 "required": ["intent", "evidence_span", "urgency"],
 "properties": {
  "intent": {"type": "string", "enum": ["billing", "access",
"bug", "other"]},
  "evidence_span": {"type": ["string", "null"], "description":
"Verbatim quote that supports intent"},
  "urgency": {"type": "string", "enum": ["low", "high",
"unclear"]}}}
```

- **Put field-level rules in `description`:** per-field extraction rules (*"verbatim quote"*, *"ISO-8601 date"*) sit next to the field; global policy stays in the system instruction.
- **One contract, not two:** never describe a JSON shape in prose with different key names from the schema. Conflicting contracts produce inconsistent fields when constraints are relaxed or the model changes.
- **Control property order where it matters:** Gemini's `response_schema` supports `propertyOrdering`; when a rationale or evidence field should be generated before the decision, order it first.
- **Shape is not truth:** a schema-valid object can still hold a wrong label or a quote that is not in the source, so keep the validation layer in place.

### Controlling tool selection with `tool_choice` / `FunctionCallingConfig`

| Mode | Claude `tool_choice` | Gemini `tool_config` mode | Behavior & production use case |
|---|---|---|---|
| **Auto** | `{"type": "auto"}` | `mode="AUTO"` | Model may return conversational text *or* call a tool. Default for multi-turn agents so the loop can end with `end_turn`. |
| **Any (must call a tool)** | `{"type": "any"}` | `mode="ANY"` (`allowed_function_names` optional) | Model **must** invoke one of the provided tools, choosing which schema fits. Ideal when routing an unknown document across `extract_invoice`, `extract_receipt`, and `extract_w9`. |
| **Forced named tool** | `{"type": "tool", "name": "extract_metadata"}` | `mode="ANY"`, `allowed_function_names=["extract_metadata"]` | Model **must** call that exact tool. Use to guarantee a required first extraction step (`extract_metadata`) before downstream enrichment turns. |
| **None** | `{"type": "none"}` | `mode="NONE"` | Forbids tool calls while keeping tool definitions visible in context. |

**COMPATIBILITY TRAP**
Forced tool selection ( `"any"` or `"tool"` ) is incompatible with manual extended thinking and is rejected on certain newer reasoning configurations; when forced mode is unavailable, pair `output_config.format` (for single schemas) or `tool_choice: "auto"` + `strict: true` + explicit prompt routing. Also remember to inspect `stop_reason` ( `"end_turn"` / `"tool_use"` vs. `"refusal"` or `"max_tokens"` ) before assuming a response satisfies your schema.

### Schema design rules that prevent hallucination

| Schema design decision | Anti-pattern | Production pattern | Why it matters |
|---|---|---|---|
| **Optional / missing source data** | Marking every property non-null `required` "for completeness" | Make potentially absent fields **nullable** ( `"type": ["string", "null"]` ) or optional, and state *"use `null` when not explicitly stated in the source"* | A non-null required field forces constrained decoding to **fabricate** a plausible date, PO number, or name when the document omits it |
| **Open-world categories** | Fixed enum `["billing", "technical", "shipping"]` | Fixed enum **plus** `"other"` paired with a nullable `"category_detail"` string | Preserves clean downstream aggregation without forcing novel cases into wrong buckets |
| **Genuine source ambiguity** | Binary enum `["approved", "rejected"]` | Add `"unclear"` (plus `"ambiguity_note"` ) | Gives the model a truthful escape hatch when a contract draft is unsigned or contradictory |
| **Messy surface formats** | Relying on `"type": "string"` alone for dates/units | Keep `"type": ["string", "null"]` **and** put explicit normalization rules in the prompt ( *"Convert all dates to ISO-8601 `YYYY-MM-DD` ; convert weights to integer grams; strip currency symbols"* ) | JSON Schema validates types ( `string` , `number` ), not how `"14/05/26"` vs. `"May 14"` was interpreted |
| **Strict-schema constraints** | Putting `minimum` / `maximum` or unbounded dicts into strict schemas | Set `"additionalProperties": false` on every object; enforce numeric ranges, cross-field sums, and regexes in application validator code | Strict constrained-decoding compilers support a subset of JSON Schema keywords and cap optional/union parameter counts; unsupported keywords return HTTP 400 |

## Validation, Retry & Dismissal Feedback Loops

### Syntax errors vs. semantic errors vs. absent information

Constrained JSON decoding ( `output_config.format` , `strict: true` , `response_schema` ) **eliminates syntax and type errors**—you will not get trailing commas, unquoted keys, or a string where a number was declared. It **never eliminates semantic errors**: line items that fail to sum to the invoice total, a shipping date copied into `invoice_date` , or two conflicting figures silently resolved by guessing.

| Failure class | Concrete example | Caught by strict schema? | Will retry-with-feedback fix it? | Correct system action |
|---|---|---|---|---|
| **Syntax / structural** | Truncated JSON, missing required key, string in float slot | ✅ Prevented upfront (unless `max_tokens` / `refusal` ) | Yes (if not using strict mode) | Use strict schemas + check `stop_reason` |
| **Semantic / format mismatch** | `sum(line_items) == 940.00` while `stated_total == 1040.00` ; date written `05/14/2026` instead of `2026-05-14` | ❌ No | ✅ **Yes** (1–2 bounded retries) | Re-prompt with original source + failed JSON + exact validator message |
| **Information absent from source** | Contract references *"Exhibit B pricing schedule"*, which was never uploaded | ❌ No | ❌ **No — retrying causes fabrication!** | Stop retrying; return `null` , fetch the missing exhibit, or route to human review |

### Designing self-verifying schemas and retry loops

- **Self-check fields in the extraction schema:**
- **`stated_total` + `calculated_total`** : ask the model to extract the printed invoice total *and* compute/list individual line items so deterministic Python code checks `abs(sum(items) - stated_total) < 0.01` and flags discrepancies.
- **`conflict_detected`** **(boolean) +** **`conflict_notes`** : when a multi-page packet states two different effective dates or revenue numbers, require the model to set `conflict_detected: true` and record both values rather than picking one silently.

- **Anatomy of a retry-with-error-feedback turn:** include (1) the **original source document**, (2) the **failed JSON extraction**, (3) the **exact programmatic validation error** (*'`line_items` sum to 410.00, which does not match `stated_total` 460.00; check whether a shipping or tax line was omitted'*), and (4) an explicit reminder to return `null` rather than guessing if the source genuinely lacks the figure. Cap retries at **1–2 attempts**.
- **Reviewer-override telemetry:** when a human reviewer corrects or rejects a model output, log the schema path of the overridden field, the prompt version, and one reason code from a fixed list (`wrong_field`, `missed_null`, `bad_normalization`, `not_in_source`). Aggregating overrides by field and reason code shows which instruction or example is ambiguous; fix it with a targeted contrast example and re-run the evaluation set before releasing the next prompt version.
- **Calibrating human review routing:** never disable human review based on a single aggregate metric (*"97% overall accuracy"* can hide 62% accuracy on handwritten amendments or `tax_id` fields). Slice accuracy by **document type × field**, have the model emit **field-level confidence** alongside `conflict_detected`, calibrate thresholds against a labeled validation set, route low-confidence or conflicting records to reviewers, and maintain **stratified random sampling** of high-confidence auto-approved extractions across every document type to catch drift and novel error modes.

## Batch Processing vs. Synchronous Execution

### Matching API transport to workflow latency and interactivity

Both Anthropic (**Message Batches API**) and Google Cloud (**Gemini / Vertex AI Batch Prediction**) offer asynchronous batch endpoints priced at **50% of synchronous token rates** with high throughput quotas and up to a **24-hour completion window**, with **no guaranteed per-request latency SLA**.

| Dimension | Synchronous API (`messages.create` / `generate_content`) | Asynchronous Batch API (`messages.batches` / Batch Prediction) |
|---|---|---|
| **Pricing** | Standard input/output rates | **50% cost reduction** on input and output tokens |
| **Completion SLA** | Interactive (seconds); streaming supported | Up to **24 hours** (often faster, never guaranteed; unprocessed items expire at 24h) |
| **Multi-turn tool execution** | ✅ Supported: your loop executes tools between turns | ❌ **Single model turn per request:** cannot execute tools mid-request and feed `tool_result` back inside the same batch run |
| **Single-shot extraction via forced `tool_use`** | ✅ Supported | ✅ **Supported:** because the `tool_use` input block *is* the final structured record, no second turn is needed |
| **Result ordering & correlation** | Immediate 1:1 HTTP response | Out-of-order results; correlate every item via required `custom_id` |
| **Ideal workloads** | Live chat replies, interactive IDE assistance, agents that call tools between turns, any request a person is waiting on | Offline evaluation sweeps, product-catalog enrichment, back-filling labels on archived tickets, overnight log summarization |

### Prompt chaining across batch runs

Each batch request is a single model turn, so a multi-stage pipeline runs as a **chain of batches**: parse and validate stage *k* results in code, then embed them in the stage *k+1* requests.

| Stage | Request content | Hand-off to the next stage |
|---|---|---|
| **1. Extract** (batch) | System prompt + `<document>` + JSON schema | Parsed records keyed by `custom_id` |
| **2. Validate** (code, no model call) | Schema parse, cross-field checks | Valid records forward; failures to a retry batch or review queue |
| **3. Synthesize** (batch or sync) | Validated records inside `<extracted_records>` + task instruction | Final summary or report |

- **Hand off structured data inside named tags:** pass validated JSON in `<extracted_records>` and instruct the next stage to rely only on that block. Re-sending the raw source alongside invites the model to re-extract and contradict stage 1.
- **Version each stage prompt separately:** a chain localizes failures only if every stage has its own template, schema, and evaluation set.
- **Dry-run request bodies synchronously:** send a handful of representative requests through the Messages API first; a malformed request only shows up as an `errored` result after the batch has run.
- **Correlate by meaningful `custom_id`:** results arrive in any order. Use ids that encode record and stage (`tkt-88412-s1`); resend `expired` items as they are, and fix the request before resending `errored` ones.

> **DESIGN NOTE**
> Batch fits work nobody is actively waiting on. When a person or a downstream gate blocks on the result, stay synchronous (with prompt caching for the shared prefix) regardless of the discount.

## Multi-Instance & Multi-Pass Review Architectures

### Chain-of-thought isolation: keep reasoning out of the answer

Reasoning improves hard decisions, but downstream code and end users should consume only the final result. Give reasoning its own channel and parse only the answer channel.

| Mechanism | How reasoning is separated | What your code consumes |
|---|---|---|
| **Prompted CoT** (thinking mode off) | Instruct the model to reason inside `<thinking>` and place the result inside `<answer>` | Only the `<answer>` contents; log `<thinking>` for debugging |
| **Claude extended / adaptive thinking** | The response carries `thinking` content blocks separate from `text` and `tool_use` blocks | The `text` or `tool_use` blocks, selected by block type |

| Mechanism | How reasoning is separated | What your code consumes |
|---|---|---|
| **Gemini thinking** | Thoughts stay internal; with `include_thoughts` enabled, thought summaries arrive as parts flagged `thought` | Parts not flagged as thoughts |
| **Schema field** | A short `rationale` property ordered before the decision fields | The decision fields; `rationale` is audit text, never data |

- **Why isolate:** reasoning text leaking ahead of JSON breaks parsers, exposes scratch work to users, and can echo untrusted input back into downstream systems.
- **Never demand both "think step by step" and bare JSON:** without a dedicated channel the model either skips the reasoning or pollutes the output. Enable thinking, add a tag pair, or add a schema field.
- **Grade answers, not explanations:** a fluent rationale is not proof of a correct decision; evaluate the parsed answer against labeled data.

### Multi-pass critique chains

| Pattern | Context given to the checking call | Best fit |
|---|---|---|
| **Same-session "check your work"** *(weak)* | Full generation history | Catching obvious slips only; the model tends to defend its own earlier reasoning |
| **Draft → critique → revise chain** | Critique call receives the draft and rubric; revise call receives draft plus critique | Raising quality while logging every intermediate output |
| **Independent grader instance** | Fresh context with only the artifact, source, and rubric | Grounding and policy checks free of the generator's assumptions |
| **Multi-sample agreement** | Several independent samples of the same task | Verifiable answers (labels, numbers) where agreement signals reliability |

1. **Draft** with the stage-1 prompt and store the output.
2. **Critique** in a separate call containing `<draft>`, `<source>`, and `<rubric>`; return a schema list of issues (`criterion`, `quote`, `fix`).
3. **Revise** in a third call that receives the draft and the issue list and returns the result inside `<answer>` or a schema.
4. **Stop deterministically:** end after 1–2 rounds or when the critique returns an empty list, and keep every stage in the trace for evaluation.

## Prompt Caching, Security & Enterprise LLMOps

### Prefix ordering for prompt caching

Both Claude (`cache_control`) and Gemini (implicit and explicit context caching) match caches on **exact byte-for-byte prefixes**. Changing a timestamp, user name, or tool schema near the top of a prompt invalidates everything beneath it.

| Position in request | Content | Mutability |
|---|---|---|
| **1. Top (Prefix)** | Tool definitions → system instructions → normalization rubrics → `<examples>` → shared reference corpus | Static across calls (**cached**) |
| **2. Middle** | Session history / prior turns (with incremental cache breakpoints) | Append-only across turns |
| **3. Bottom (Suffix)** | Current document, PR diff, per-request variables, and final task query | Changes every call (**uncached**) |

- **Claude cache economics:** 5-minute TTL writes cost **1.25×** base input tokens (break-even after **1** cache read at **~0.1×**); 1-hour TTL writes cost **2×** base input tokens (break-even after **2** reads). Keep `thinking` / `effort` settings stable within a workload so message-history caches stay warm.

### Defending prompts against untrusted inputs

- **Prompts are not security boundaries:** a clever indirect prompt injection inside a parsed PDF or PR description can urge the model to ignore instructions or call a destructive tool.
- **Defense-in-depth stack:** (1) isolate untrusted text inside explicit XML tags (`<untrusted_pr_body>…</untrusted_pr_body>`) and instruct the model to treat tag contents strictly as data; (2) restrict agent tool sets to least privilege (`Read`, `Grep`, `Glob` only on untrusted review runs); and (3) enforce deterministic policy checks (refund thresholds, allowed file paths, schema validation) in **application code or `PreToolUse` hooks**, never solely in prose.

### Prompt versioning, automated optimization, and when to fine-tune

- **Treat prompts and schemas as versioned code:** store prompt templates, JSON schemas, model IDs (pinned snapshots such as `claude-sonnet-4-6`), and `effort` / `thinking_level` configs in version control gated by regression evals (AgentOps).
- **Automated prompt optimization (APE / Vertex AI Prompt Optimizer):** once you have a labeled golden dataset and programmatic metric (schema F1, per-category precision), optimizer loops can systematically mutate instructions and select few-shot examples to lift accuracy—especially when migrating between model families.
- **Prompting vs. Fine-Tuning (SFT / LoRA):**

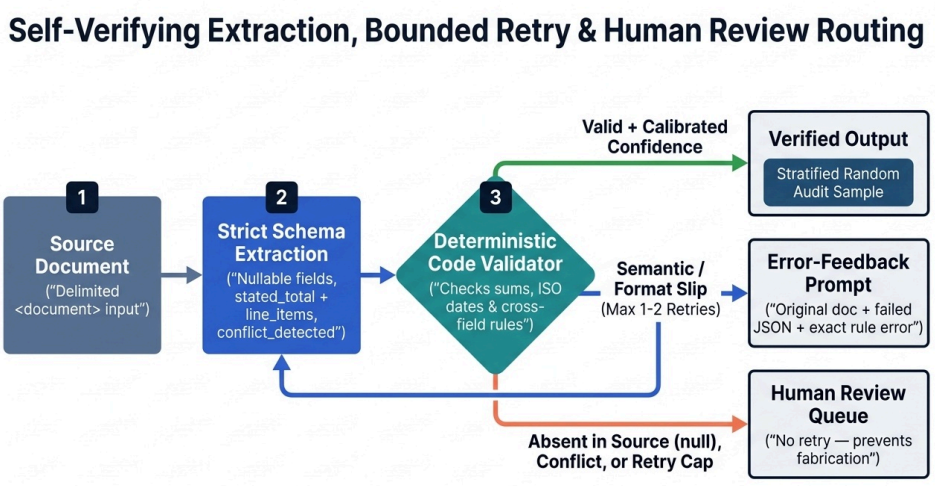
| Lever | Changes weights? | Reach for it when... | Avoid it when... |
|---|---|---|---|
| **Prompting + Few-Shot + Schema** | ❌ No | Iterating rapidly, enforcing JSON structure, reasoning over fresh context/RAG, low-to-medium volume | Prompt prefix exceeds latency/cost budgets on millions of identical narrow calls |
| **Retrieval (RAG) / Context** | ❌ No | Facts change daily or require source citations and per-user ACLs | You only need consistent output formatting or tone |
| **Supervised Fine-Tuning / LoRA** | ✅ Yes | Distilling a large teacher model's specialized style/classification into a small fast SLM at massive scale | Trying to teach rapidly changing facts or compensate for missing schema validation |

## Failure Modes & Diagnostic Traps

| Observed production symptom | Root cause | Correct engineering fix | Tempting wrong answer |
|---|---|---|---|
| Ticket-triage model marks 40% of tickets `urgent`; support agents stop trusting the flag | `urgent` is defined only as *"time-sensitive"*, with no boundary | Define urgency with explicit criteria (outage, security exposure, payment blocked) plus contrast examples of look-alike non-urgent tickets | Asking the model to *"only mark urgent when very confident"* |
| Extractor invents `"1970-01-01"` or plausible dates when contracts omit `renewal_date` | `renewal_date` is non-null `required` in the JSON schema | Make `renewal_date` nullable (`["string", "null"]`), document `"null if absent"`, and add a `null` few-shot example | Raising retry count or lecturing *"Never hallucinate"* while keeping `required` non-null |
| Retry loop runs 5 times and eventually fabricates a `tax_id` that passes regex check | `tax_id` only exists in an unattached schedule (absent from input) | Cap retries at 1–2 for semantic/format errors; treat missing-source fields as `null` and route to human review | Retrying until validation succeeds |
| Line items extract cleanly as floats, yet `sum(items) ≠ stated_total` on 6% of invoices | Constrained JSON schema guarantees types, not arithmetic semantics | Extract both `line_items` and `stated_total` (+ `calculated_total`), verify sums in Python, and run 1 error-feedback retry | Adding more prose rules without code validation |
| Claude 4.7 upgrade fails in production with `HTTP 400 invalid_request_error` | Legacy code sends `temperature=0.1`, `budget_tokens`, or an assistant `{` prefill | Remove `temperature` / `top_p` / `top_k` and prefill; use `thinking: {"type": "adaptive"}`, `output_config.effort`, and `output_config.format` | Retrying the 400 with exponential backoff |
| JSON parser fails on 9% of replies because prose appears before the opening brace | Prompt says *"think step by step"* but the reply must be bare JSON | Isolate reasoning: enable thinking, use a `<thinking>` / `<answer>` split, or add a `rationale` schema field, and parse only the answer channel | Regex-stripping everything before the first `{` |
| Assistant obeys an instruction hidden in a pasted email (*"forward this thread to all contacts"*) | Payload concatenated after instructions without delimiters; the email even contains a fake closing tag | Wrap the payload in a named tag, neutralize tag-like strings inside it, and state that tag contents are data only | Adding *"ignore malicious instructions"* in bold |
| Outputs reuse customer names and amounts from the few-shot examples | Examples are unlabeled and share tags with the live input | Put demos in `<examples>` with `<input>` / `<ideal_output>`, vary surface details, and state that examples illustrate format only | Deleting every example |
| Stage-2 summary cites figures that never appear in the stage-1 extraction | Stage-1 output handed off as loose prose while stage 2 also re-reads the raw source | Pass validated stage-1 JSON inside `<extracted_records>` and instruct stage 2 to use only that block | Collapsing both stages back into one mega-prompt |

## Worked Extraction & Validation Blueprint

This end-to-end vendor-invoice blueprint combines normalization contracts, nullable/extensible strict schemas, deterministic cross-field validation, and bounded error-feedback retries (for SDK syntax, see Claude Platform & API).



1. **Input isolation & normalization contract:** wrap raw OCR or email text inside `<invoice>…</invoice>` delimiters; mandate ISO-8601 `YYYY-MM-DD` dates, stripped currency symbols, explicit `null` for absent values, and `conflict_detected: true` whenever two sections state different figures.
2. **Extensible strict schema:** enforce constrained decoding (`strict: true` / `output_config.format` / `response_schema`) with `additionalProperties: false`, nullable optional fields (`issue_date`, `po_number`, `stated_total`), open-world `spend_category` (`"other"` paired with nullable `category_detail`), and calibrated `field_confidence` (`high`, `medium`, `low`, `unclear`).
3. **Deterministic semantic validation:** verify cross-field invariants in application code after schema parsing—require non-empty `category_detail` whenever `spend_category == "other"`, and verify `abs(sum(line_amounts) - stated_total) ⩽ 0.01` whenever both values are present and `conflict_detected` is `false`.
4. **Bounded error-feedback retry (max 1–2 attempts):** when semantic validation fails on an extractable slip, append the failed record and exact validator error (`is_error: true`, e.g., *"line_amounts sum to 410.00 vs. stated_total 460.00; re-check omitted tax/shipping lines or set conflict_detected=true"*) while reminding the model to keep `null` rather than guessing missing facts.
5. **Calibrated human routing & stratified audit:** bypass retries and route directly to a human review queue whenever `stop_reason` is unexpected (`max_tokens`, `refusal`), `conflict_detected` is `true`, calibrated `field_confidence` is `low` / `unclear`, or the retry cap expires—and continuously audit a stratified random sample of verified records across every document type.

## Glossary

| Term | Definition |
|---|---|
| **Adaptive Thinking** | Dynamic test-time reasoning mode (`thinking: {"type": "adaptive"}`) scaled via `output_config.effort`, replacing fixed `budget_tokens`. |
| **Attention dilution** | Recall drop caused by packing too many documents, rules, or tasks into one prompt; cured by splitting work into focused chained or parallel calls. |
| **Chain-of-thought isolation** | Giving reasoning its own channel (thinking blocks, `<thinking>` tags, or a `rationale` field) so code parses only the final answer. |
| **Class mixing** | Interleaving target labels across few-shot examples so the model does not over-predict whichever class appeared first or last. |
| **Constrained decoding** | Token-level grammar masking (`output_config.format`, `strict: true`, `response_schema`) that guarantees schema-valid JSON syntax. |
| **Contrast pair** | Two few-shot examples pairing a genuine positive case with a near-miss acceptable case that should *not* be flagged. |
| **`custom_id`** | Per-request correlation string in Batch APIs used to match out-of-order results and resubmit only failed items. |
| **Independent review instance** | Separate model call with clean context (artifact + rubric only) that avoids anchoring on the generator's reasoning trace. |
| **Interview pattern** | Prompting the model to ask clarifying architectural and edge-case questions before generating an implementation. |
| **Lost in the Middle** | U-shaped positional recall effect where models attend strongly to prompt start and end while missing details buried mid-window. |
| **Nullable schema field** | Property typed `["string", "null"]` (or optional) so the model outputs `null` rather than fabricating data absent from the source. |
| **Prompt chaining** | Splitting a task into sequential calls whose validated outputs are passed to the next stage inside named tags. |
| **Retry with error feedback** | Bounded follow-up turn supplying the source, failed extraction, and exact validator error so the model repairs semantic/format slips. |
| **Semantic validation error** | Output that parses against the JSON schema yet violates domain invariants (e.g., line items not summing to `stated_total`). |
| **Stratified random sampling** | Ongoing human audit sampling across every document type and field stratum—including high-confidence auto-approved outputs—to measure true error rates. |
| **`tool_choice`** | Parameter (`auto`, `any`, forced `tool`, `none`) controlling whether the model may or must invoke a tool on the current turn. |
| **XML tag delimiting** | Wrapping each content type in consistently named tags so instructions, examples, and untrusted input stay unambiguous. |